

OpenSource RISC Processors: An introduction into OpenRISC and RISC-V

Advanced System-on-Chip Design

Lecture 6


05.04.2016

Michael Gautschi

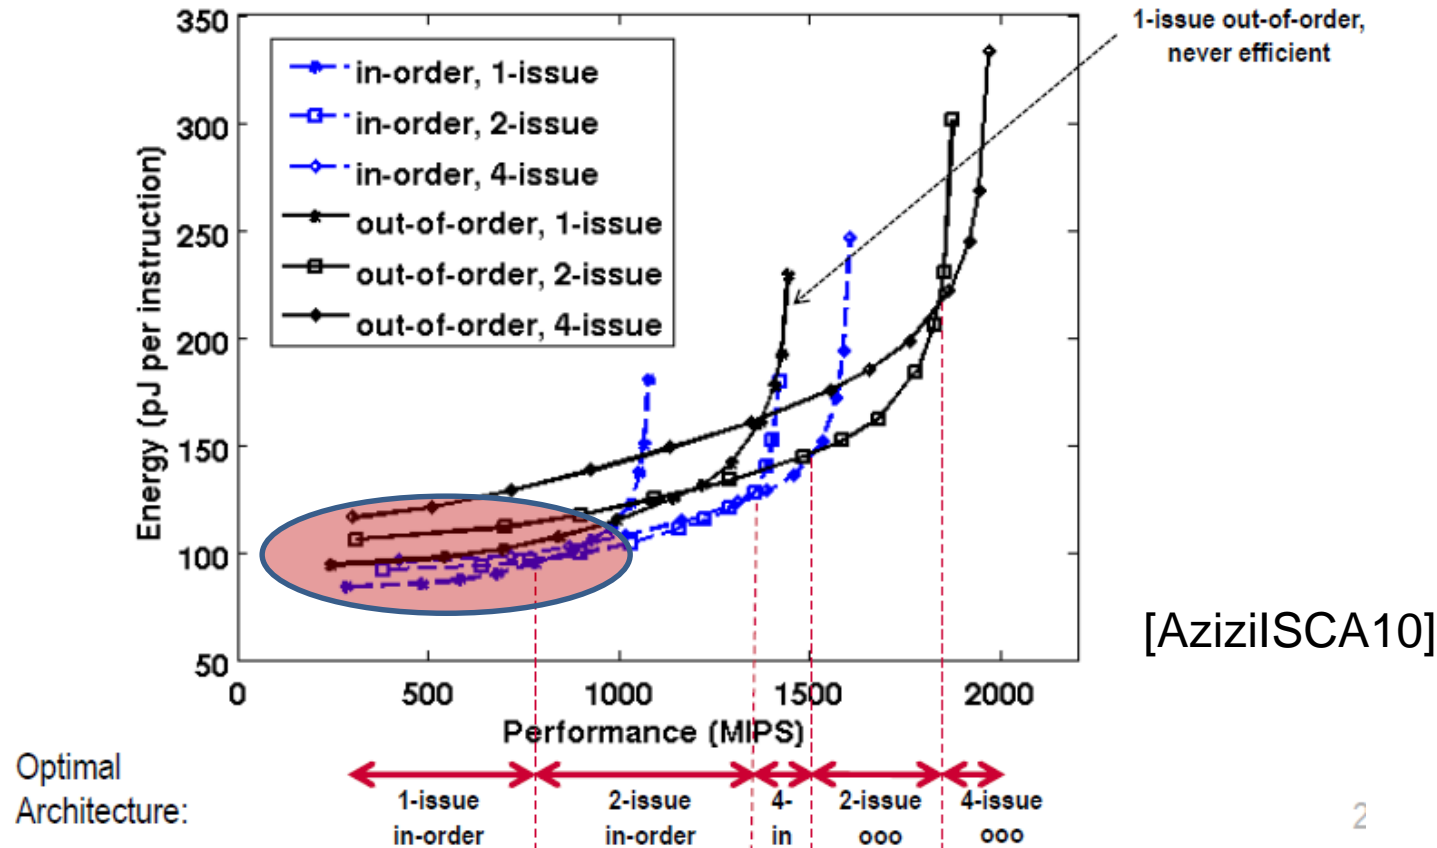
IIS-ETHZ

Prof. Luca Benini

IIS-ETHZ



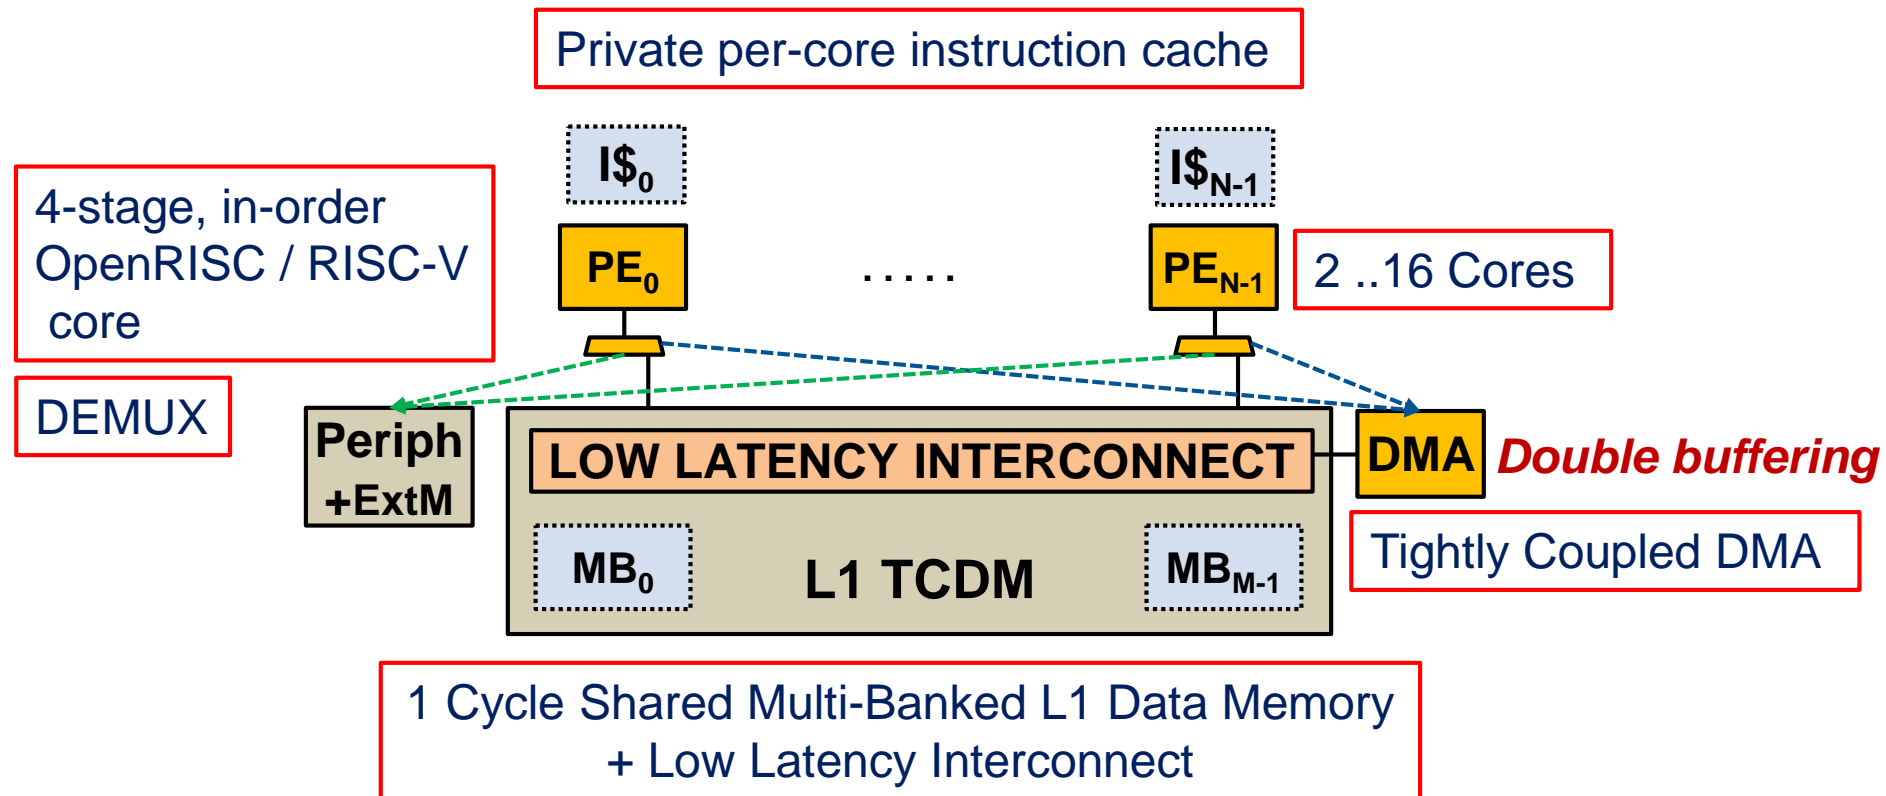
Introduction – the “best” core



- Single issue in-order is most energy efficient
- Put more than one + shared memory to fill cluster area

Introduction – Building PULP

SIMD + MIMD + sequential



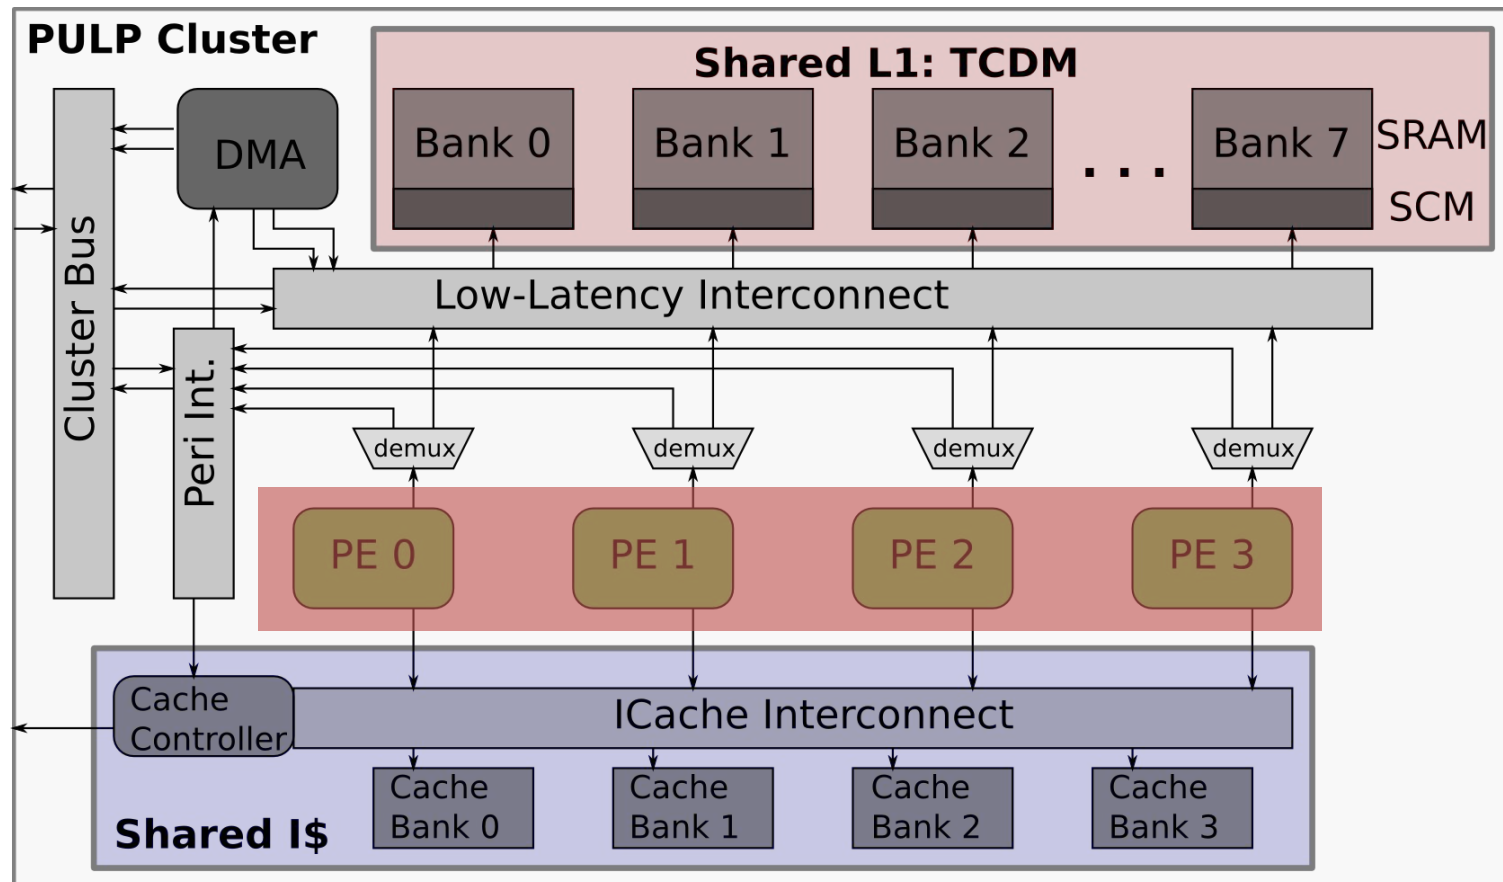
“GPU like” shared memory → low overhead data sharing

Near Threshold but parallel → Maximum Energy efficiency when Active

+ strong power management for (partial) idleness

Introduction – PULP Architecture

- Uses Open Source RISC processor
 - OpenRISC, RISC-V ISA

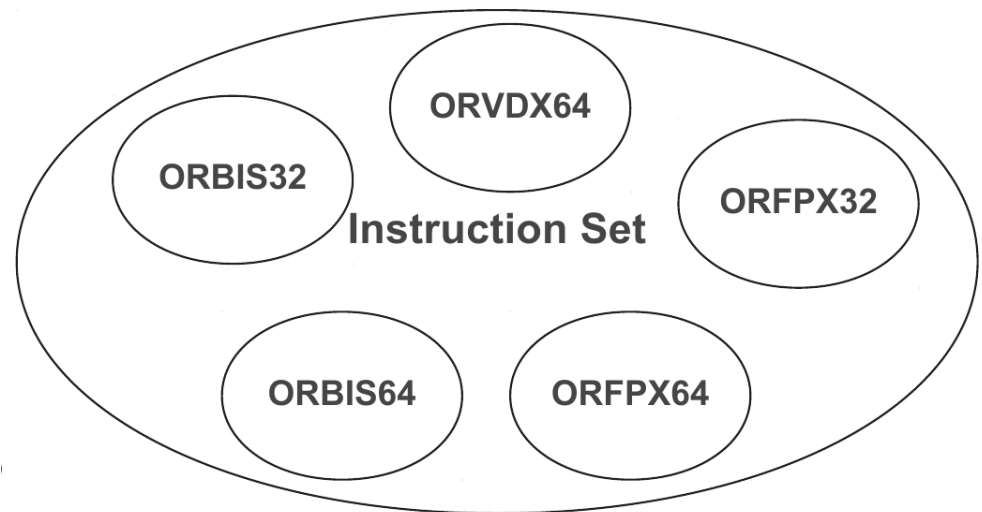


Introduction

- Outline:
 - OpenRISC Instruction-set
 - Basic instruction set
 - Micro-architecture
 - Organization of the pipeline
 - Instruction set extensions for improved performance
 - Hardware and software impact
 - RISC-V architecture
 - Difference to OpenRISC
 - Exercise session about OpenRISC/ RISC-V processor cores
 - Exercise session
- Goals:
 - Learn how to run applications on the Pulpino architecture
 - Understand the impact of the presented hardware extensions
 - Including some pro/ and cons

OpenRISC Instruction Set

- Open source 32-/64bit RISC architecture
 - Similar to MIPS architecture described in Hennessey/Patterson
 - ORBIS32:
 - 32-bit integer instruction
 - 32-bit load/store instructions
 - Program flow instructions
 - ORBIS64:
 - 64-bit integer instructions
 - 64-bit load/store instructions
 - ORFPX32:
 - Single precision floating point instru
 - ORFPX64:
 - Double-precision floating point instructions
 - ORVDX64:
 - 64-bit vector instructions
- ⇒ In the following we focus on the 32-bit ORBIS32 instruction set!




OpenRISC Instruction Set

- ORBISX32 consists of three types of instructions
 - R-type instructions:
 - Register - register operations
 - Examples:
 - ALU operations: *l.add, l.mul, l.sub, l.or, l.mac, etc.*
 - Comparisons: *l.sfeq, l.sfges, etc.*

l.add

Add

l.add



l.sfeq

Set Flag if Equal


l.sfeq




OpenRISC Instruction Set

- ORBISX32 consists of three types of instructions
 - I-type instructions:
 - Operations with an immediate
 - Examples:
 - Load/store operations: *l.lwz, l.sw, l.lhz, l.lbz etc.*
 - ALU operations: *l.addi, l.muli, l.ori, etc.*
 - Comparisons: *l.sfeqi, l.sfnei. etc.*

l.lwz Load Single Word and Extend with Zero l.lwz



l.muli Multiply Immediate Signed l.muli



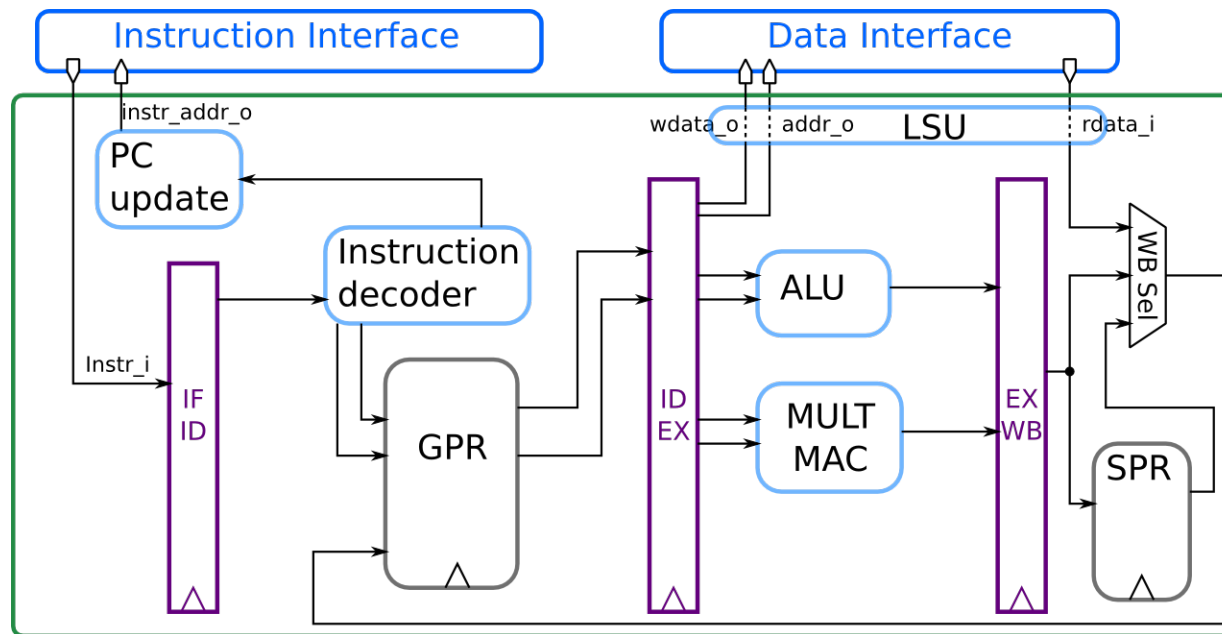
- ORBISX32 consists of three types of instructions
 - J-type instructions:
 - Jumps and branches
 - Examples:
 - Jump instructions: *l.j, l.jal, l.rfe, etc.*
 - Conditional branches: *l.bf, l.bnf*

l.jal



OpenRISC Micro-architecture:

- Core architecture which has been originally developed here at IIS in a semester thesis.
 - The architecture is called OR10N
 - It has been improved over the years and become a good core architecture
- Simple four-stage pipeline architecture: IF, ID, EX, WB
- Single cycle memory access





Register file & special purpose registers(SPR)

- Register file organization:
 - 32 registers 32-bit registers
 - Most important registers are:
 - r0 = always zero
 - r1 = stack pointer
 - r9 = link register, holds function return address
 - r11/r12 = return values
- Special purpose registers:
 - Status register contains flags {overflow, carry, branch}
 - Contains registers which are not regularly accessed:
 - Interrupt controller configuration
 - Timer
 - Data/instruction cache control
 - Debug unit
 - Performance counters

Load/Store Unit

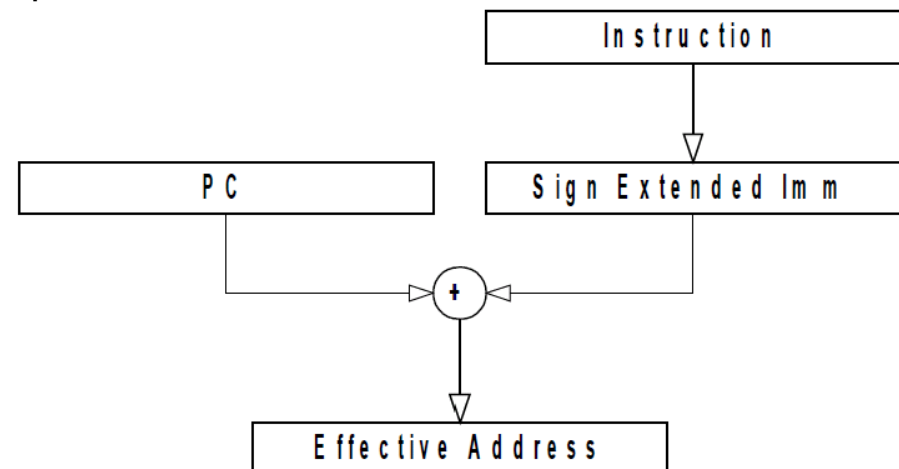
- 32 bit load-store interface
- Supported instructions:
 - Load word/halfword/ byte
 - With zero or sign extension
- Addressing mode aligned data requests
 - l.lwz/s word aligned
 - l.lhz/s half word aligned
 - l.lbz/s byte aligned
- Stall pipeline if exception has been detected
 - Access to protected address
 - Unaligned access
- No out of order requests

l.lwz Load Single Word and Extend with Zero l.lwz



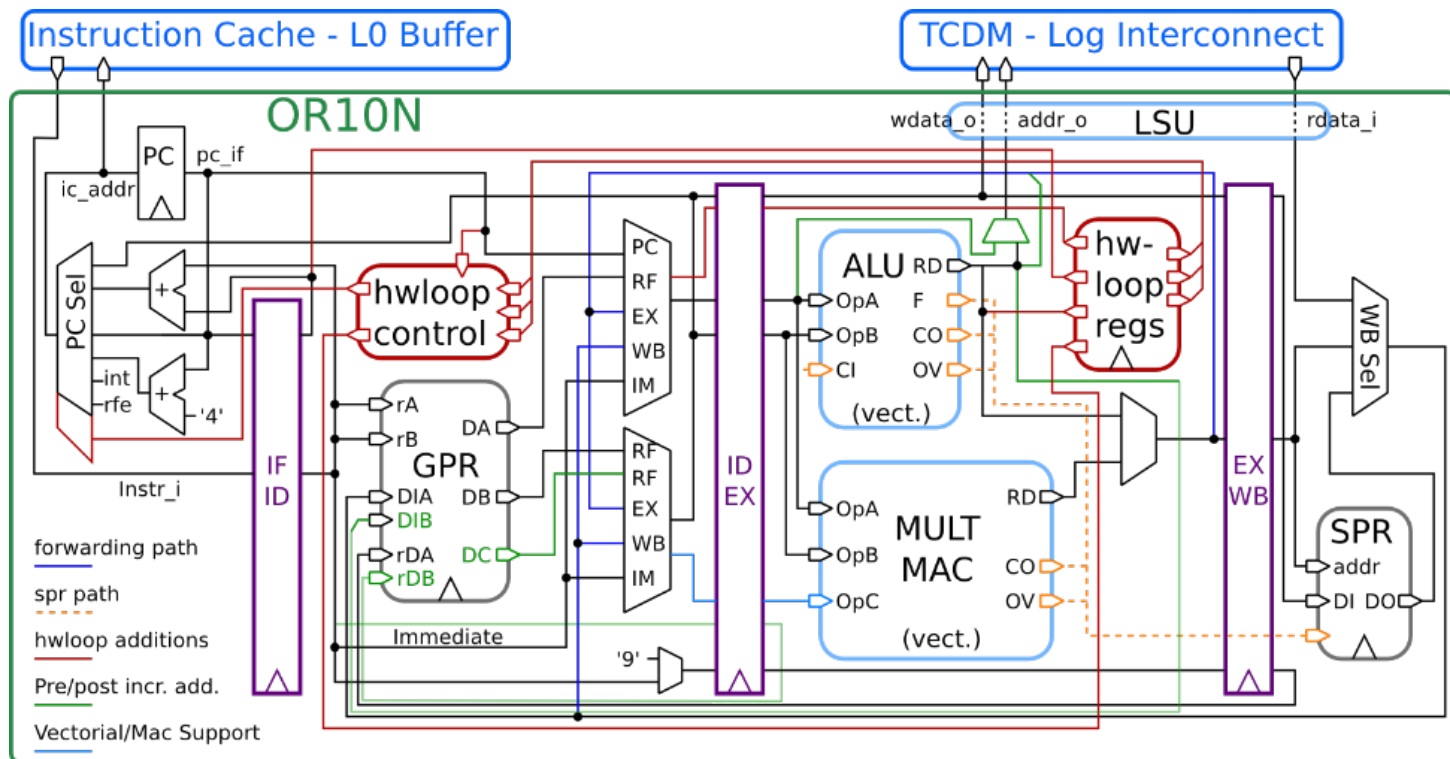
OpenRISC: Control Flow

- Branches
 - l.bnf : jump to PC + sign extended immediate if flag is not set
 - l.bf : jump to PC + sign extended immediate if flag is set
 - Delay slot is always executed
- Jumps:
 - l.jr : jump to address stored in a register
 - l.jalr : jump to address stored in a register and link r9 to instruction after delay slot
 - l.j : jump to PC + sign extended immediate
 - l.jal : jump to PC + sign extended immediate and link r9
 - l.rfe : return from exception, jump to EPCR
- No support for VLIW
 - Instructions are always 32 bit



OR10N Instruction Extensions for OR10N:

- In order to improve performance and efficiency of the core we have evaluated several instructions and added the following instructions:
 - Hardware loops
 - Pre/post memory address update
 - New MAC
 - Vector unit
 - Unaligned memory access



Instruction Extensions: Hardware Loops

- Hardware loops or Zero Overhead Loops can be implemented to remove the branch overhead in for loops.
- After configuration with start, end, count variables no more comparison and branches are required.
- Smaller loop benefit more!
- Loop needs to be set up beforehand and is fully defined by:
 - Start address
 - End address
 - Counter

```

1c0006f8: 00 68 24 e2 e2246800 l.add r17, r4, r13
1c0006fc: 00 00 6f 8e 8e6f0000 l.lbz r19, 0(r15)
1c000700: 00 00 31 8e 8e310000 l.lbz r17, 0(r17)
1c000704: 06 8b 33 e2 e2338b06 l.mul r17, r19, r17
1c000708: 01 00 ad 9d 9dad0001 l.addi r13, r13, 1
1c00070c: 00 60 91 e1 e1916000 l.add r12, r17, r12
1c000710: 20 00 0d bc bc0d0020 l.sfeqi r13, 32
1c000714: f9 ff ff 0f 0fffffff9 l.bnf -0x1c # 0x1c0006f8
1c000718: 20 00 ef 9d 9def0020 l.addi r15, r15, 32

```

9 loop instructions

Hardware loop support

```

1c0006ec: 20 00 20 09 09200020 lp.counti L1, 0x20
1c0006f0: 00 00 60 a9 a9600000 l.ori r11, r0, 0x0
1c0006f4: 04 00 20 08 08200004 lp.start L1, 0x10
1c0006f8: 09 00 a0 08 08a00009 lp.end L1, 0x24
1c0006fc: 00 00 80 a9 a9800000 l.ori r12, r0, 0x0
1c000700: 00 00 a7 a9 a9a70000 l.ori r13, r7, 0x0
1c000704: 00 60 e4 e1 e1e46000 l.add r15, r4, r12
1c000708: 00 00 2d 8e 8e2d0000 l.lbz r17, 0(r13)
1c00070c: 00 00 ef 8d 8def0000 l.lbz r15, 0(r15)
1c000710: 06 7b f1 e1 e1f17b06 l.mul r15, r17, r15
1c000714: 01 00 8c 9d 9d8c0001 l.addi r12, r12, 1
1c000718: 00 58 6f e1 e16f5800 l.add r11, r15, r11
1c00071c: 20 00 ad 9d 9dad0020 l.addi r13, r13, 32

```

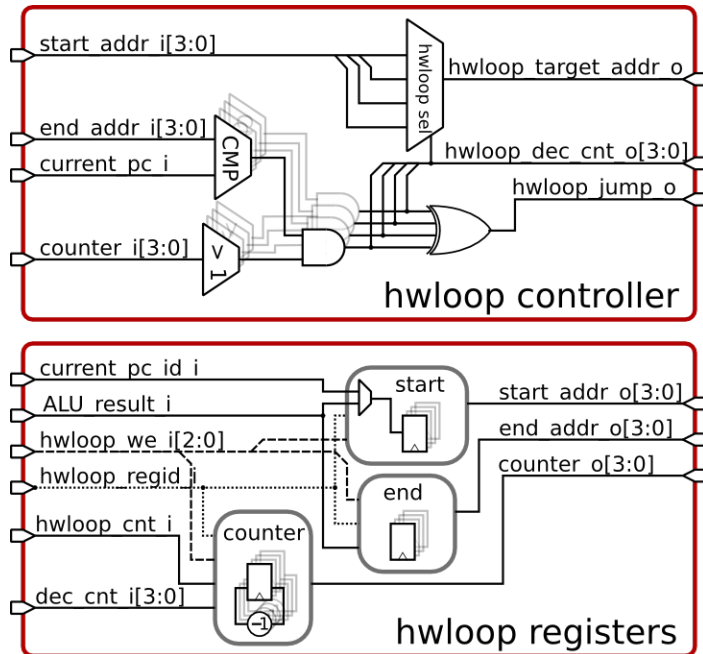
3 setup instructions +

7 loop instructions

Instruction Extensions: Hardware Loops

- Two sets registers implemented to support nested loops.
- Area costs:
 - Processor core area increases by 5%
- Performance:
 - Speedup can be up to factor 2!

- Hardware loop setup with:
 - 3 separate instructions
 $lp.start$, $lp.end$, $lp.count$,
 $lp.counti$
 \Rightarrow No restriction on start/end address
 - Fast setup instructions
 $lp.setup$, $lp.setupi$
 \Rightarrow Start address= PC + 4
 \Rightarrow End address= start address + offset
 \Rightarrow Counter from immediate/register



Instruction format and Opcode	Semantics
$lp.start\ J, S$ (eg. $lp.start\ L0, 10$) 000010 000 JJ SSSSSSSSSSSSSSSSSSSSS	$HWLP_START[J] = sext(S*4) + PC$
$lp.end\ J, S$ (eg. $lp.end\ L0, -8$) 000010 001 JJ EEEEEEEEEEEEEEEEEEEEE	$HWLP_END[J] = sext(E*4) + PC$
$lp.counti\ J, C$ (eg. $lp.counti\ L0, 8$) 000010 010 JJ CCCCCCCCCCCCCCCCCCCC	$HWLP_COUNT[J] = zext(C)$
$lp.count\ J, rA$ (eg. $lp.count\ L0, r5$) 000010 011 JJ AAAAA-----	$HWLP_COUNT[J] = [rA]$
$lp.setupi\ J, E, C$ (eg. $lp.setupi\ L0, 4, 8$) 000010 100 JJ CCCCCCCCCCCCCCEEEEEEE	$HWLP_START[J] = PC + 4$ $HWLP_END[J] = zext(E*4) + PC$ $HWLP_COUNT[J] = zext(C)$
$lp.setup\ J, E, rA$ (eg. $lp.setup\ L0, 8, r5$) 000010 101 JJ AAAAAEEEEEEEEEEEEEEEEEE	$HWLP_START[J] = PC + 4$ $HWLP_END[J] = zext(E*4) + PC$ $HWLP_COUNT[J] = [rA]$

Instruction Extensions: Post increment

- Automatic address update
 - Update base register with computed address after the memory access
 - ⇒ Save instructions to update address register
 - Post-increment:
 - Base address serves as memory address
- Offset can be stored in:
 - Register
 - Immediate

1c0006ec:	20 00 20 09	09200020	lp.counti	L1, 0x20	}	3
1c0006f0:	00 00 60 a9	a9600000	l.ori	r11, r0, 0x0		
1c0006f4:	04 00 20 08	08200004	lp.start	L1, 0x10		
1c0006f8:	09 00 a0 08	08a00009	lp.end	L1, 0x24	}	7
1c0006fc:	00 00 80 a9	a9800000	l.ori	r12, r0, 0x0		
1c000700:	00 00 a7 a9	a9a70000	l.ori	r13, r7, 0x0		
1c000704:	00 60 e4 e1	e1e46000	l.add	r15, r4, r12		
1c000708:	00 00 2d 8e	8e2d0000	l.lbz	r17, 0(r13)		
1c00070c:	00 00 ef 8d	8def0000	l.lbz	r15, 0(r15)		
1c000710:	06 7b f1 e1	e1f17b06	l.mul	r15, r17, r15		
1c000714:	01 00 8c 9d	9d8c0001	l.addi	r12, r12, 1	}	
1c000718:	00 58 6f e1	e16f5800	l.add	r11, r15, r11		
1c00071c:	20 00 ad 9d	9dad0020	l.addi	r13, r13, 32		

post increment support

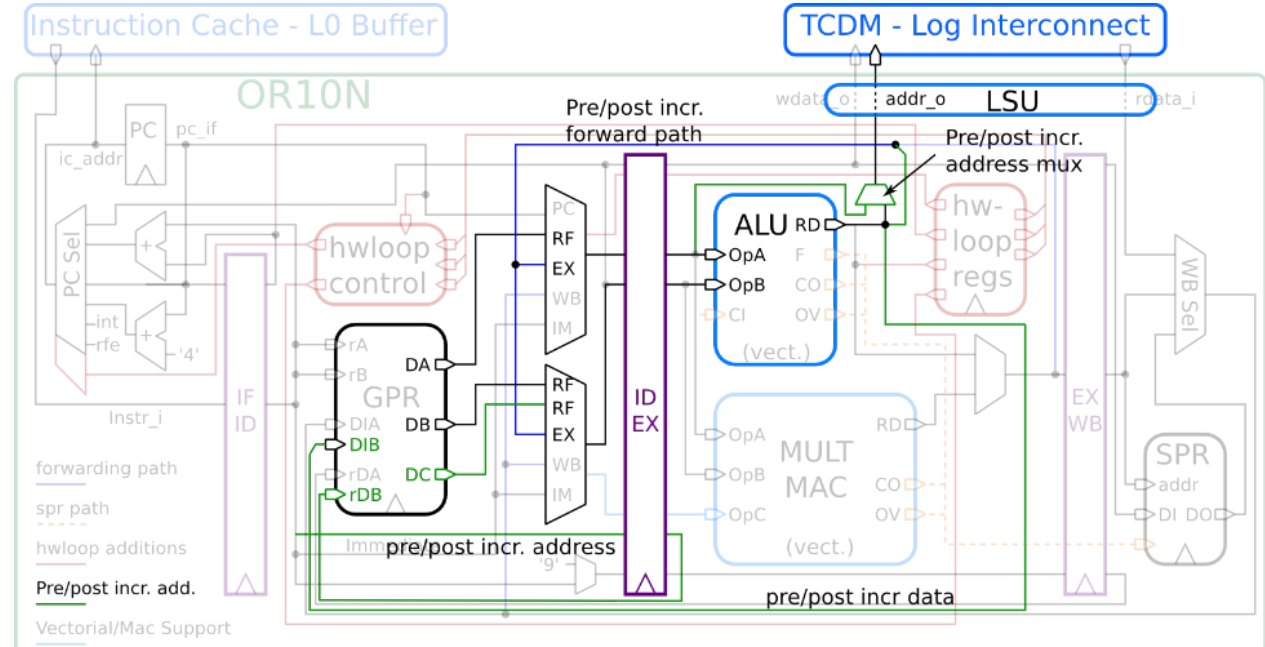
1c0006fc:	20 00 20 09	09200020	lp.counti	L1, 0x20	}	5
1c000700:	00 00 80 a9	a9800000	l.ori	r12, r0, 0x0		
1c000704:	05 00 20 08	08200005	lp.start	L1, 0x14		
1c000708:	08 00 a0 08	08a00008	lp.end	L1, 0x20		
1c00070c:	00 00 a4 a9	a9a40000	l.ori	r13, r4, 0x0		
1c000710:	00 00 eb a9	a9eb0000	l.ori	r15, r11, 0x0	}	
1c000714:	00 00 27 aa	aa270000	l.ori	r17, r7, 0x0		
1c000718:	2c 00 6d 5a	5a6d002c	l.lbz	r19, 1(r13)		
1c00071c:	0c 04 af 5a	5aaf040c	l.lbz	r21, 32(r15)		
1c000720:	06 9b 75 e2	e2759b06	l.mul	r19, r21, r19		
1c000724:	20 00 31 9e	9e310020	l.addi	r17, r17, 32	}	
1c000728:	00 60 93 e1	e1936000	l.add	r12, r19, r12		

⇒ save 2 additional instructions to track the read address of the operands!

Instruction Extensions: Post increment

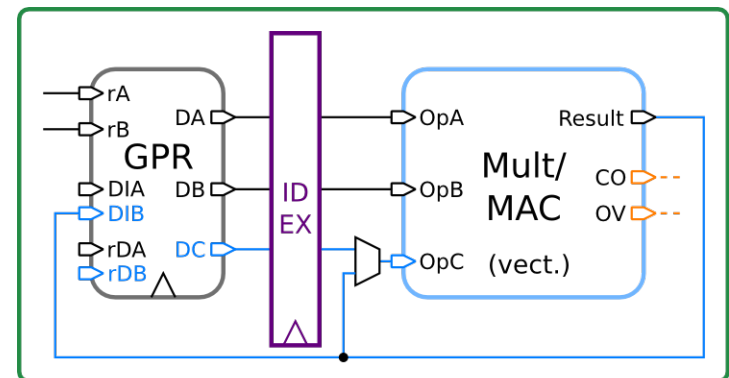
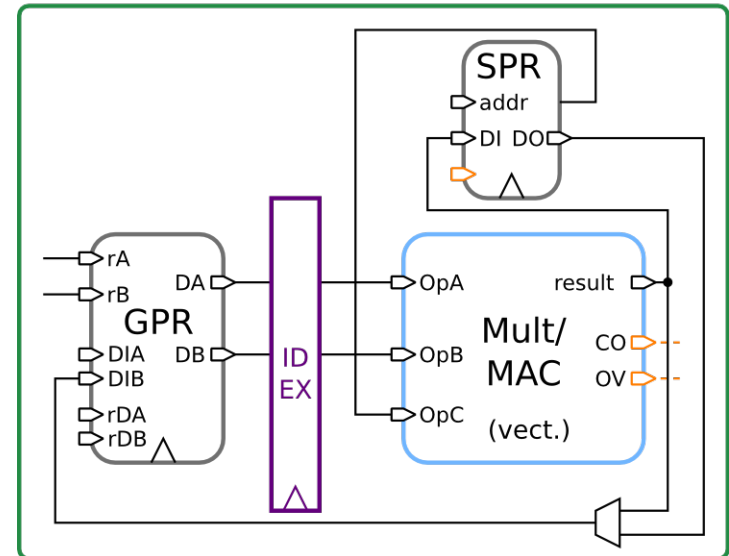
- Register file requires additional write port
- Register file requires additional read port if offset is stored in register
- Processor core area increases by 8-12 %
 - Ports can be used for other instructions

Old Instruction format	Opcode
1.s[bhw] I(rA), rB	MMMMMM IIIIII AAAAA BBBB IIIIIIIIIII
1.l[bhw][zs] rD, I(rA)	MMMMMM DDDDD AAAAA IIIIIIIIIIIIIII
New Instruction format	Opcode
1.s[bhw] I(rA!), rB	010100 MMMMM AAAAA BBBB IIIIIIIIIII
1.s[bhw] I(!rA), rB	010100 MMMMM AAAAA BBBB IIIIIIIIIII
1.l[bhw][zs] rD, rC(rA)	010110 MMMMM AAAAA DDDDD CCCCC -----
1.l[bhw][zs] rD, I(rA!)	010100 MMMMM AAAAA DDDDD IIIIIIIIIII
1.l[bhw][zs] rD, I(!rA)	010100 MMMMM AAAAA DDDDD IIIIIIIIIII
1.l[bhw][zs] rD, rC(rA!)	010111 MMMMM AAAAA DDDDD CCCCC -----
1.l[bhw][zs] rD, rC(!rA)	010111 MMMMM AAAAA DDDDD CCCCC -----



New MAC: Accumulation on register file


- Old MAC:
 - Accumulation on special 64 bit register
- New MAC:
 - Accumulation only on 32 bit data
 - Directly on the register file
- Pro:
 - Faster access to mac accumulation
 - Many accumulations in parallel
 - Single cycle mult/mac
- Contra:
 - Additional read port on the register file
 - can be used for pre/post increment with register



Instruction Extensions: 32 bit Vector Support

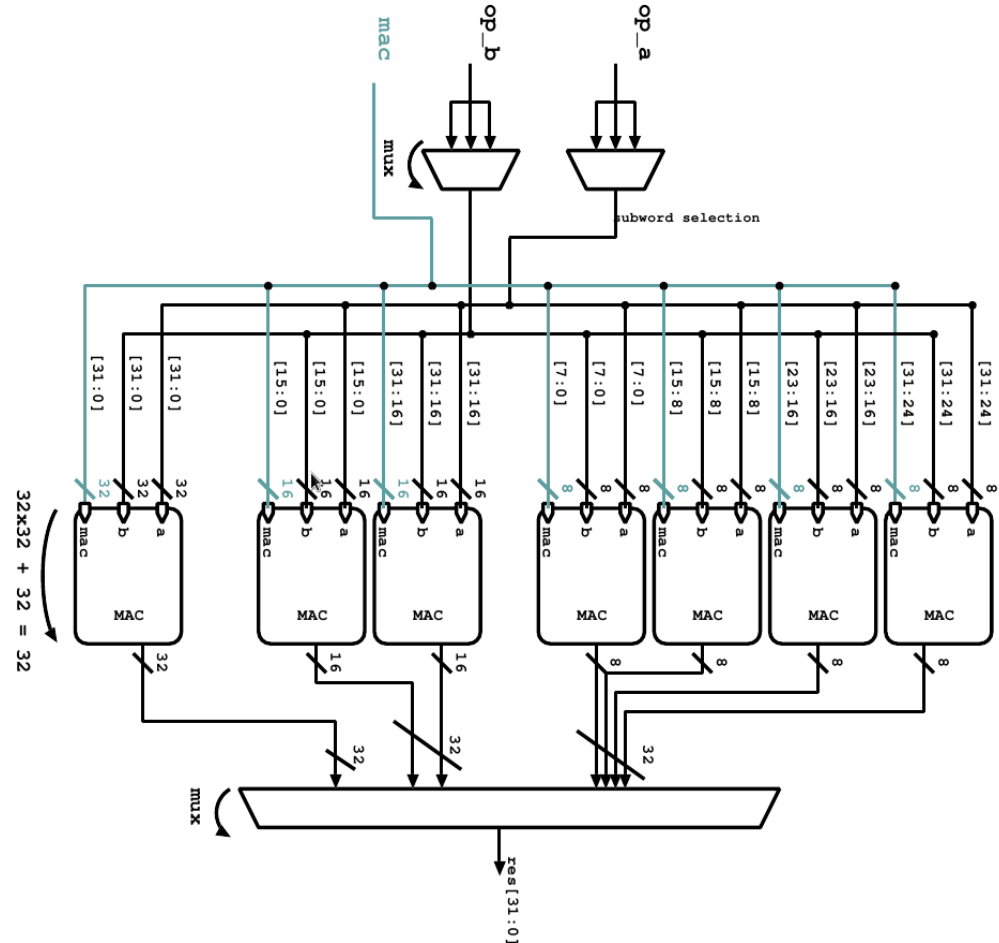
- Vector modes: (bytes, halfwords, word)
 - 4 byte operations
 - With byte select
 - 2 halfword operations
 - With halfword select
 - 1 word operation
- Vector ALU supports:
 - Vector additions
 - Vector subtractions
 - Vector comparisons:
 - Rise flag if *any*, or *all* conditions are true
- Fused vector Mult/Mac:
 - Dynamic range problem in vector multiplications

Vectorial Adder:



32 bit Vector Support: Vectorial Multiplication

- Multiplication on
 - Word,
 - Halfword
 - Byte
- 3 different multipliers:
 - 4* (8x8=8) mults.
 - 2* (16x16=16) mults.
 - 1* (32x32=32) mult.
- Output has same dynamic range as input!
 - Overflows fast in 8bit case



Performance Improvements:

Hardware loops (H) & pre-/post increment (IH):

Up to 1.75x speedup

Optimized MAC (IHM)


Up to 2.25x speedup

Vector extensions (IHVM)

Up to 5x speedup but only rarely used => Dot product

Cortex M4: STM32F429ZI

Slower, no intrinsic used



32 bit Vector Support: Dot Product Multiplier

- Dot Product: (half word example)

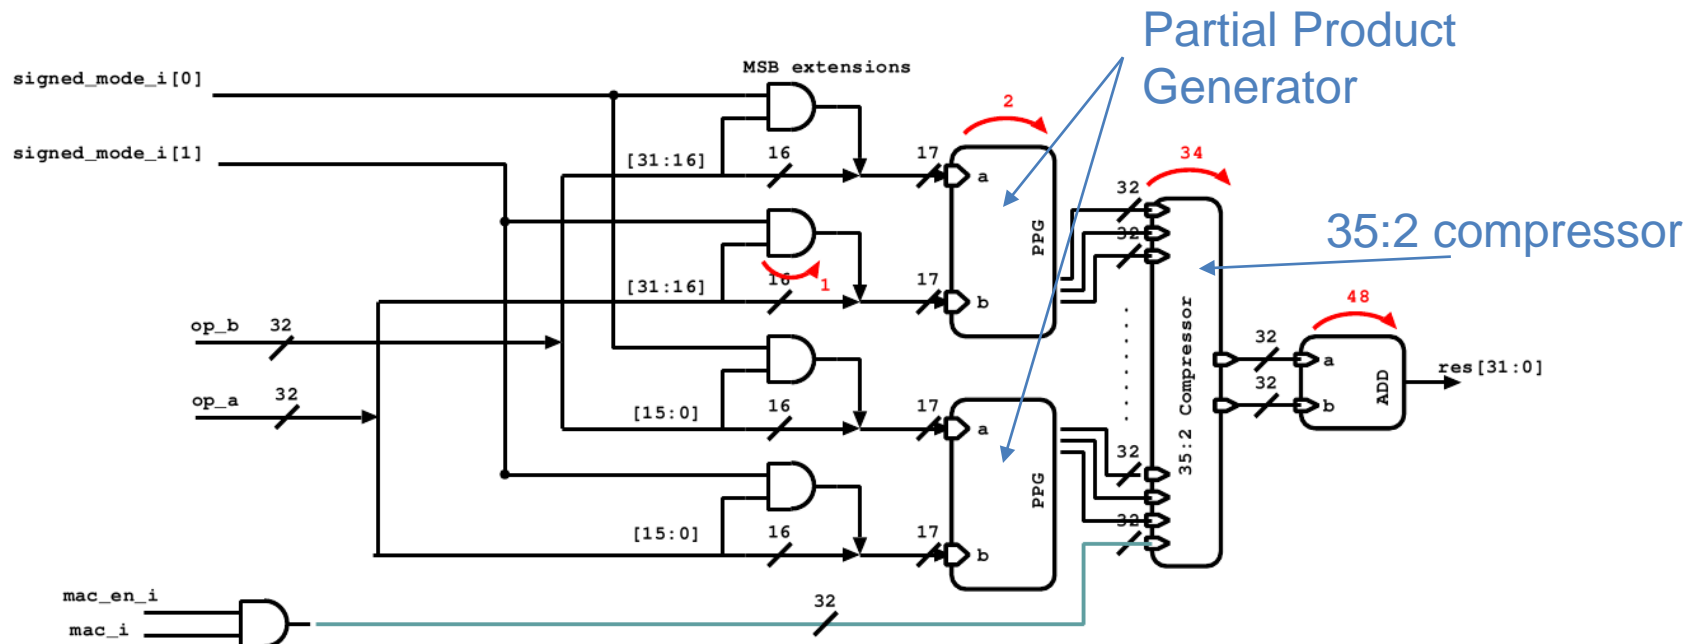
$$C[31:0] = A[31:16] \cdot B[31:16] + A[15:0] \cdot B[15:0] + C[31:0]$$

32 bit

32 bit

32 bit

=> 2 multiplications, 1 addition, 1 accumulation in 1 cycle!



Instruction Extension: Fractional Support

- Fractional format in Q-format [1]
 - Q3.12 means 1 sign bit, 3 integer bits and + 12 fractional bits (16 bit in total)

- Q3.12 format in a 32 bit reg:

-----SIIFFFFFFFFFFF

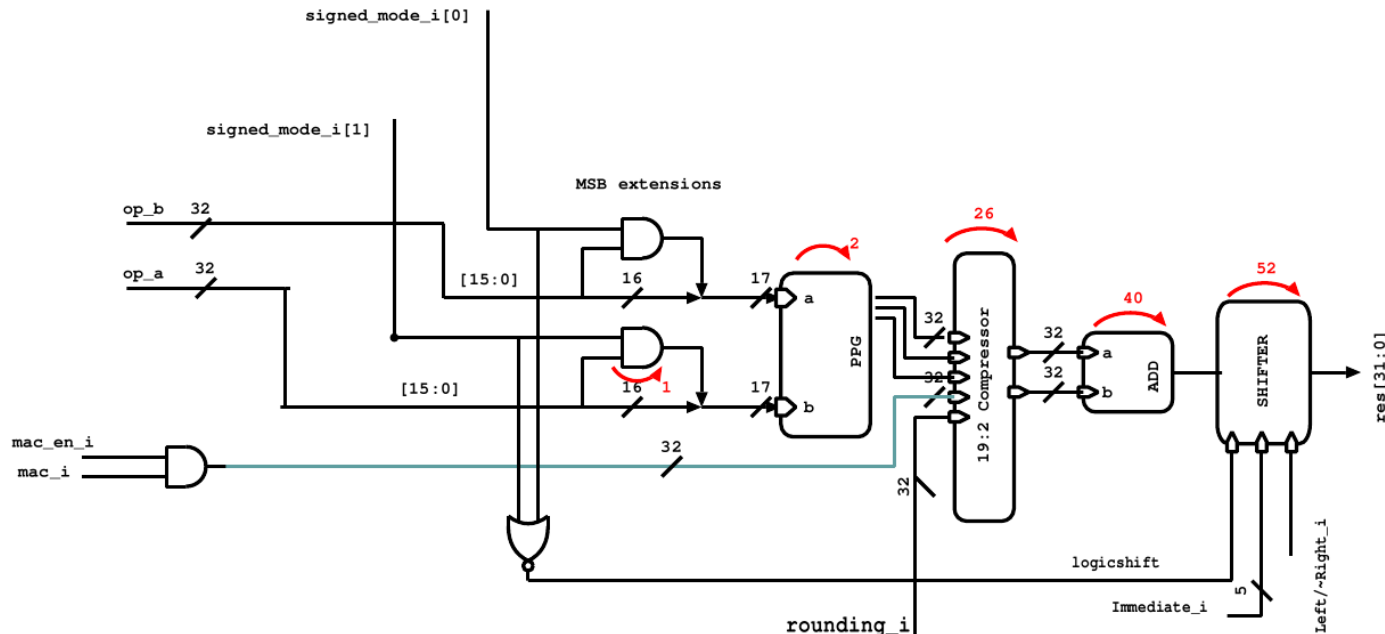
- Multiplication in Q format requires a shifter:

– $Q3.12 * Q3.12 =$

SSIIIIIIFFFFFFFFFFF

– $Q3.12 * Q3.12 \gg 12 =$

-----SSIIIIIIFFFFFFFFFFF

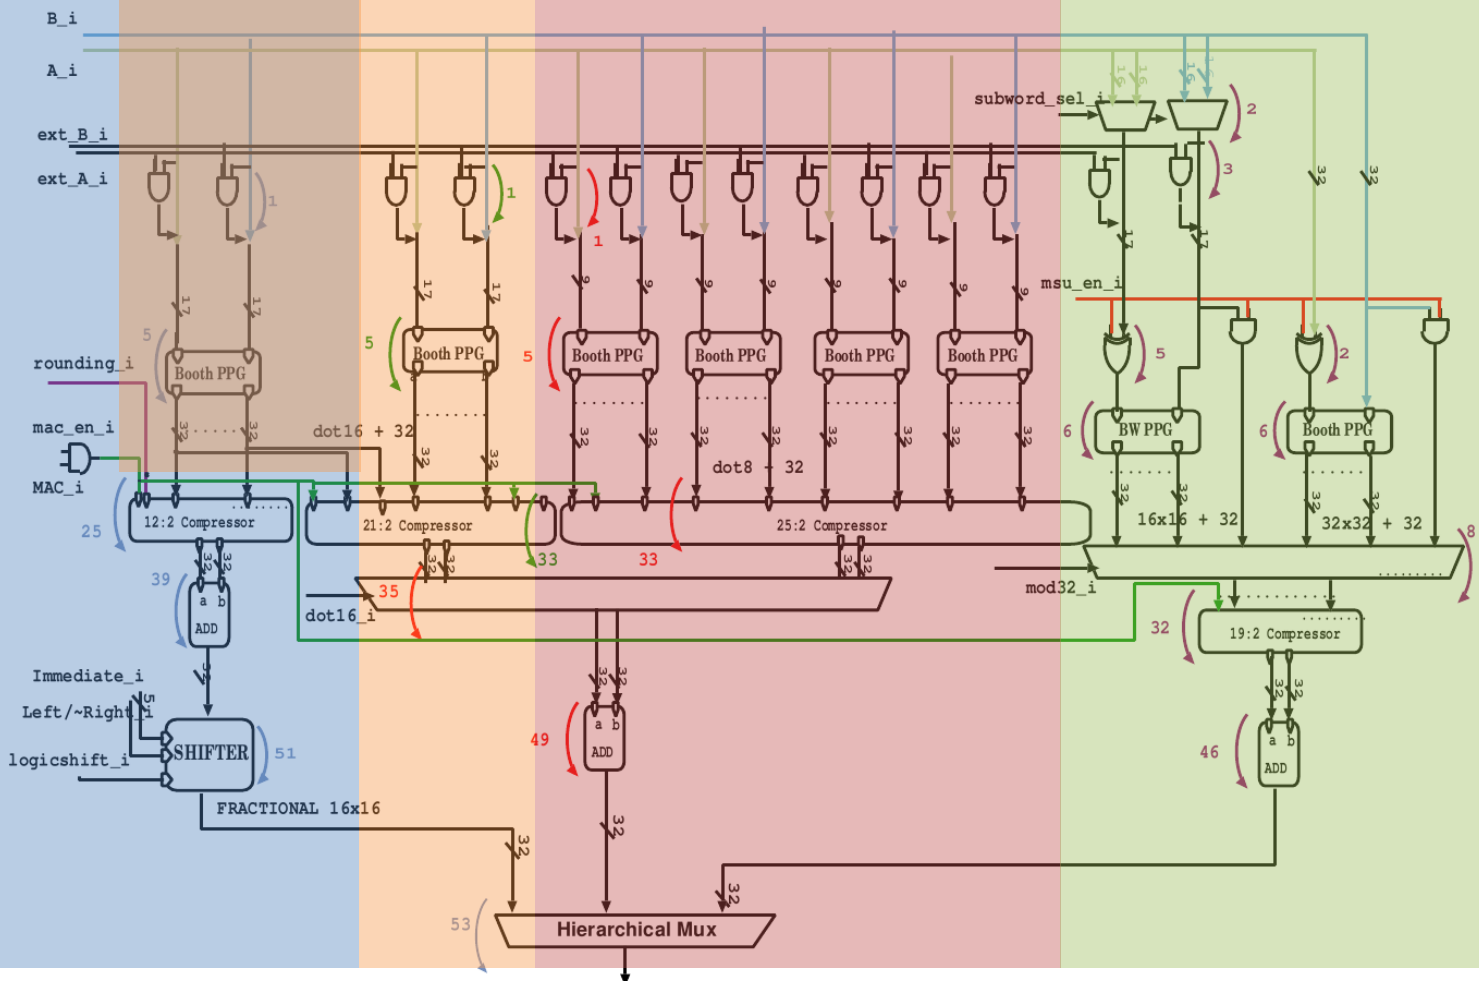


16b Dot Product Mult.

Fractional Multiplier

8b Dot Product Mult.


Integer Multiplier



32 bit Vector Support: Full Multiplier Architecture

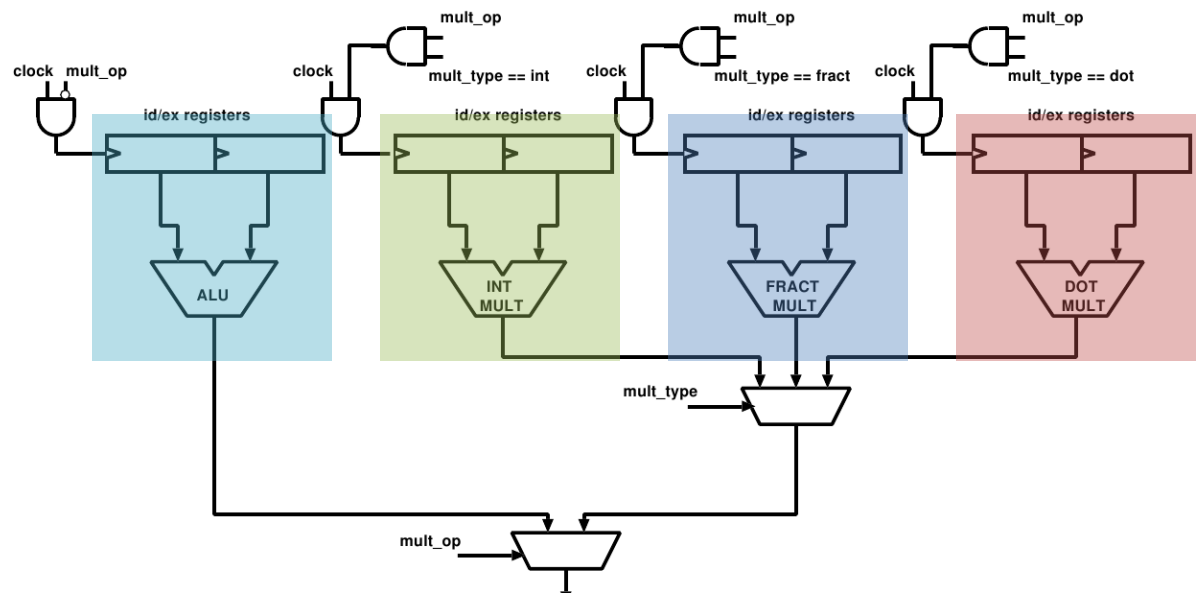

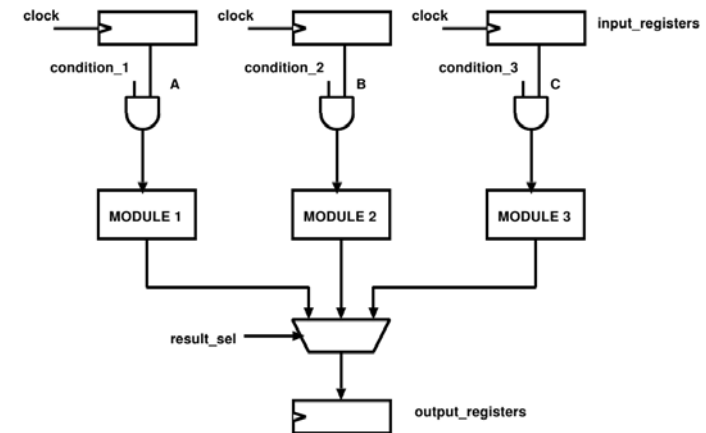
- Supports:
 - Fractional multiplications
 - Dot Products for 16b vectors
 - Dot Products for 16b vectors
 - Integer multiplication 32b*32b into 32b

- Multiplier costs:
 - Up to 2.3x bigger without resource sharing
 - Sharing makes arch. slower
 - Power ?



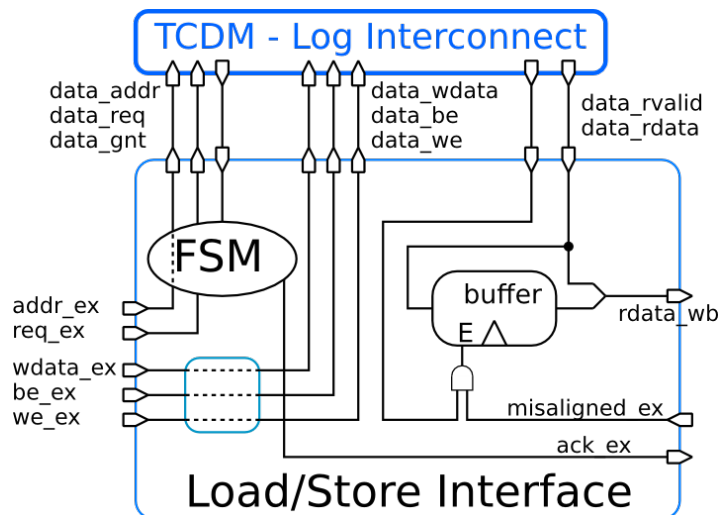
32 bit Vector Support: Power Reduction

- Operand Isolation
 - Gating inputs to zero
- Separate input registers:
 - Eliminating all active power of unused units
 - 30-60% power reduction with respect to a design without separate input registers





32 bit Vector Support: Unaligned memory access

- Unaligned memory access with 32 bit data interface:
 - Difficult to read/write unaligned words, because memories are 32 bit wide
 - Possible with multibanked memories
 - But significant hardware costs
 - Area and timing
- Implemented with two subsequent memory requests



Example: stencil with vector



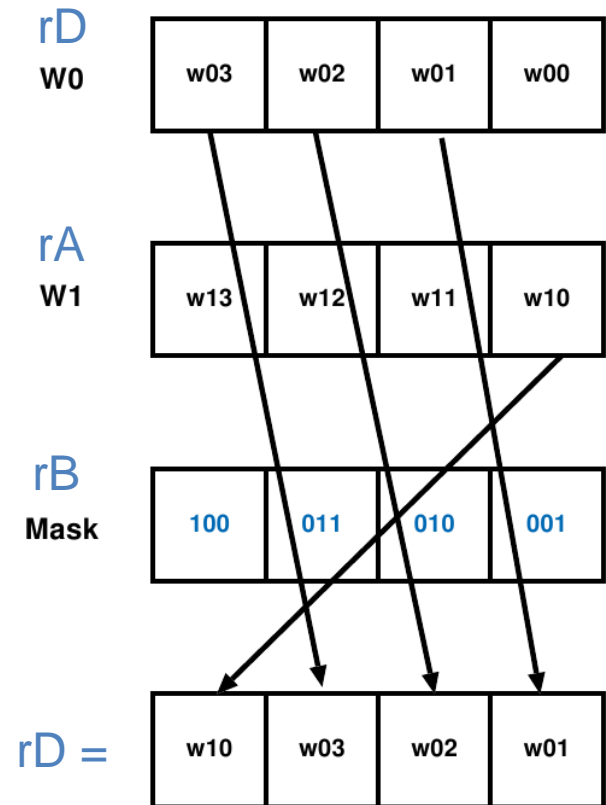
32 bit Vector Support: Shuffle Instruction

- In order to use the vector unit the elements have to be aligned in the register file
- Shuffle allows to recombine bytes into 1 register:

- lv.shuffle2.b rD, rA, rB

- $rD\{3\} = (rB[26]==0) ? rA:rD) \{rB[25:24]\}$
- $rD\{2\} = (rB[18]==0) ? rA:rD) \{rB[17:16]\}$
- $rD\{1\} = (rB[10]==0) ? rA:rD) \{rB[9:8]\}$
- $rD\{0\} = (rB[2]==0) ? rA:rD) \{rB[1:0]\}$
- With $rX\{i\} = rX[(i+1)*8-1:i*8]$

Mask bits



32 bit Vector Support: Summary

- ALU extensions for packed SIMD are easy
 - Vector addition, comparisons, etc.
- Multiplier is more complicated. Dot product is very useful for simple kernels
 - convolutions, multiplications etc.
- Load store unit needs support for unaligned access
 - Good trade-off with support in two cycles
- Vector recombination instructions are very important for the compiler!
 - To recombine bytes/ halfwords directly in registers instead of reloading everything from memory

The RISC-V ISA

- Modern ISA created by UC Berkeley for their research
- Available for 32-bit, 64-bit and 128-bit
- Little-endian
- Published as Free and Open RISC ISA
- The ISA specifications were previously controlled by UCB, now shifting to the RISC-V foundation
- RISC-V foundation is controlled by the members
Everyone can become member, just costs a bit of money
- See official website <http://riscv.org>

The RISC-V ISA: Introduction

- Generally kept very simple and extendable
- Separated into multiple specifications
 - User-Level ISA spec (compute instructions)
 - Compressed ISA spec (16-bit instructions)
 - Privileged ISA spec (supervisor-mode instructions)
 - More to come
- Implementations:
 - We have a similar architecture than in OpenRISC
 - Can be used in Semester/Master - Projects

User-Level ISA Spec

- Defines the normal instructions needed for computation
- Spec separated into “extensions”
- Defines a mandatory base integer instruction set: “I extension”
- ISA support is given by RV + word-width + extensions supported
 - E.g. **RV32I** means 32-bit RISC-V with support for the I instruction set
- I: Integer instructions; alu, branches, jumps, loads and stores
 - Support for misaligned memory access is mandatory
- M: Multiplication and (!!!) Division
- A: Atomic instructions
- F: Single-Precision Floating-Point
- D: Double-Precision Floating-Point
- C: Compressed Instructions (more later)

User-Level ISA Spec: Standard and Non-standard extensions

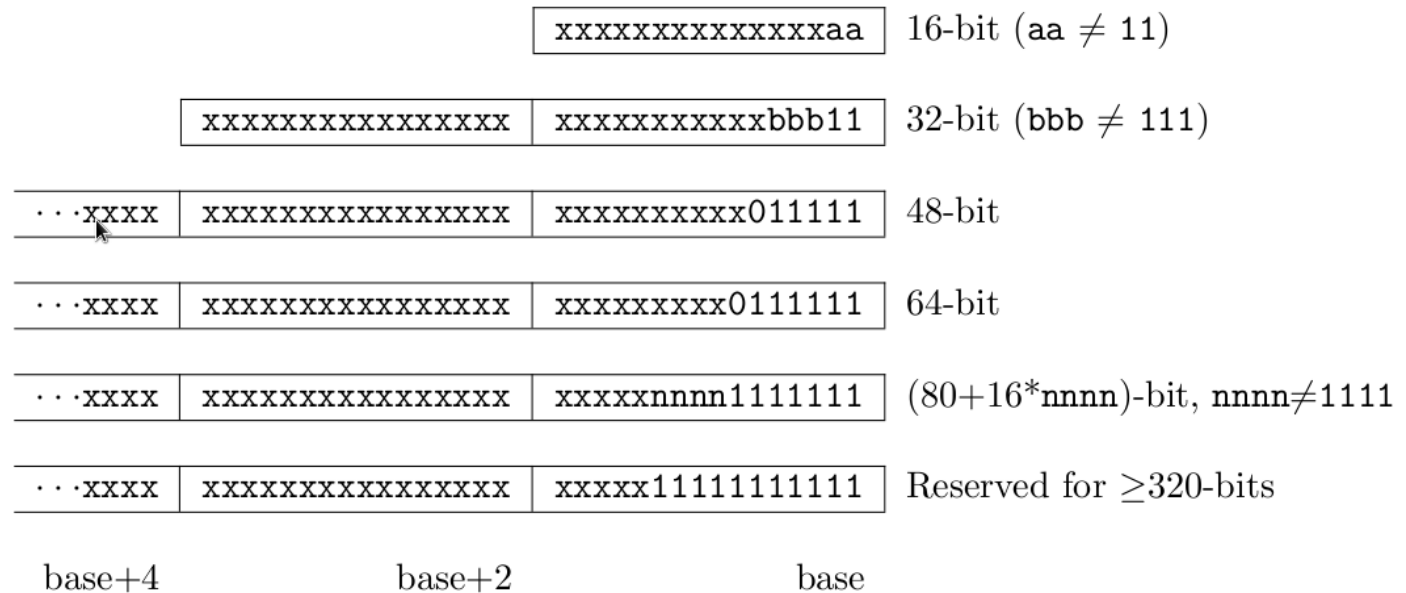
- Extensions mentioned so far are so called **Standard Extensions**
- Reserved opcodes for standard extensions
- Rest of opcodes free for non-standard implementations
- Standard extensions will be frozen and will not change in the future



011	100	101	110	111 ($> 32b$)
C-MEM	OP-IMM	AUIPC	OP-IMM-32	48b
MO	OP	LUI	OP-32	64b
ADD	OP-FP	<i>reserved</i>	<i>custom-2/rv128</i>	48b
AL	SYSTEM	<i>reserved</i>	<i>custom-3/rv128</i>	$\geq 80b$

Instruction Length Encoding

- Supports by design 16, 32, 48, 64, ... bit long instruction words




Compressed Instruction Spec (Draft)

- Still a draft, but will be frozen very soon if there are no complaints
- Compressed instructions are 16-bit wide
 - Allows to reduce code size
 - needs support for misaligned instruction memory access
- Compressed instruction ISA spec is no ISA per se
 - All compressed instructions map to I instructions, can be expanded
 - Preprocessing step for instructions needed or separate decoding for compressed instructions

Compressed Instruction Spec (Draft)

Claim code size reduction by ~34%



Differences to OpenRISC: No flags, no delay slot

- Branches:

OpenRISC

```
l.sfeq rA, rB  
l.bf 0x20  
l.nop
```

RISC-V

```
beq rA, rB, 0x20
```

- Jumps are more general

OpenRISC

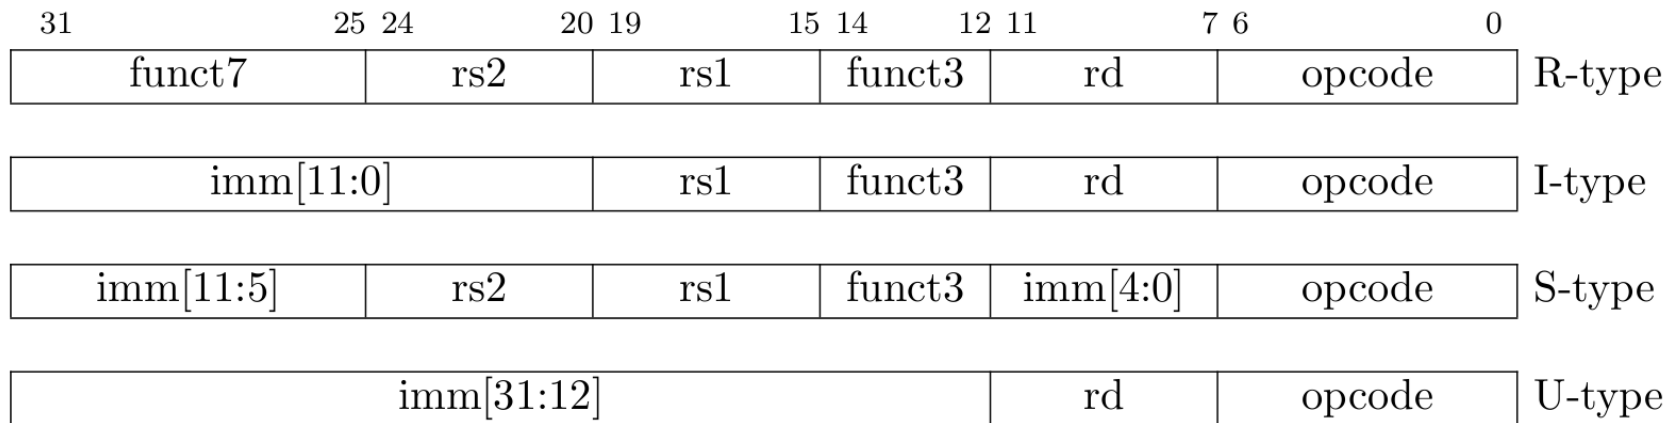
```
l.jalr rB
```

RISC-V

```
jalr rD, rs1, 0x10
```

Differences to OpenRISC: Opcode Sizes

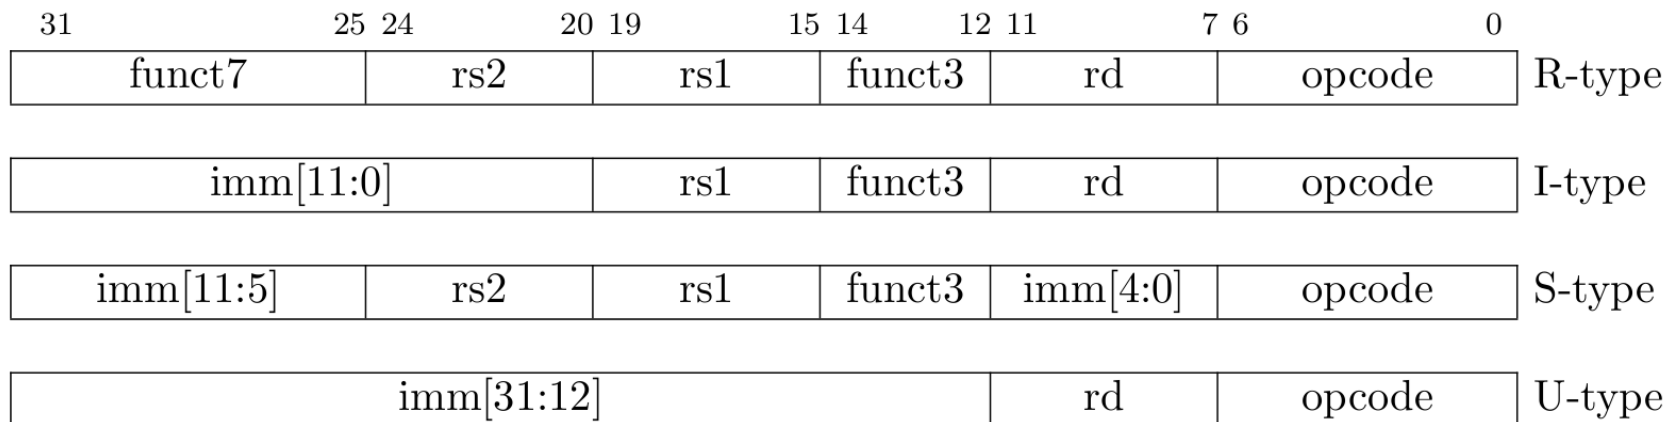
- RISC-V: 7-bit main opcodes
 - actually only 5-bits due to compressed instructions
 - sub-opcodes (funct*) for most instruction types



- OpenRISC: 6-bit main opcodes

Differences to OpenRISC: Immediate Size

- RISC-V: 12-bit immediates, all sign-extended
 - sign-bit always in same position



- OpenRISC: 16-bit immediates, mixed sign- and zero-extended

Differences to OpenRISC: Constructing a 32-bit number

- OpenRISC has 16 bit immediates with zero-extension
 - Just use 2 instructions
- RISC-V has 12 bit immediates with sign-extension
 - load upper immediate first and then add sign-extended value
 - upper immediate is set to imm+1 to correct for sign-extension if necessary

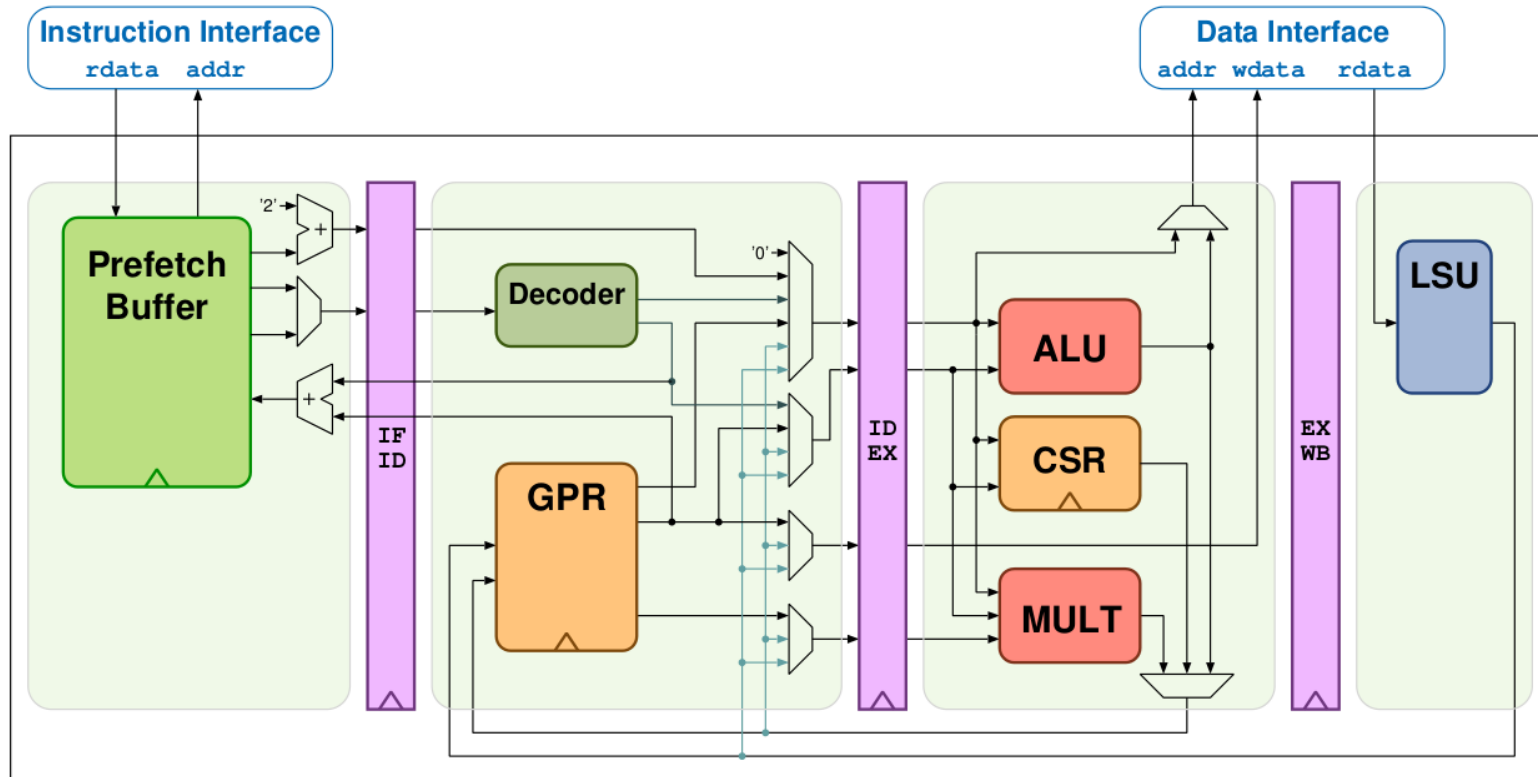
OpenRISC

```
l.movhi r1, 0x1000  
l.ori   r1, r1, 0x1800
```

RISC-V

```
lui  x1, 0x10002  
add  x1, x1, -2048
```

RI5CY: Core architecture (4 stage pipeline)

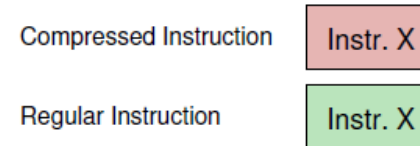


- CSR in EX instead of WB

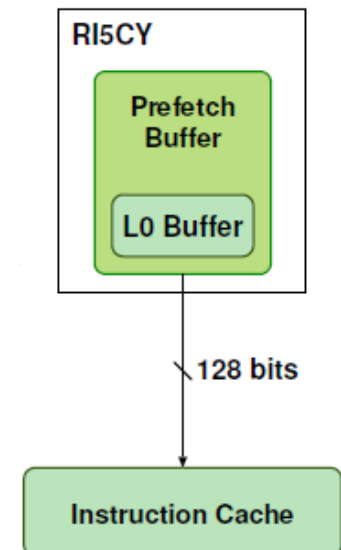
RI5CY: Prefetcher: Why is branching difficult

- No delay slot in RISC-V:
 - Jumps loose one cycle
 - Next instruction already fetching and probably ready in IF stage
- Combined branches: No setflag instruction
 - Branch decision computed with branching instruction
 - Branch decision computed in EX stage
 - Taken branches loose two cycles
 - Branch only available late in pipeline
 - Not taken branches don't loose cycles

RI5CY: Prefetcher: The need for prefetching

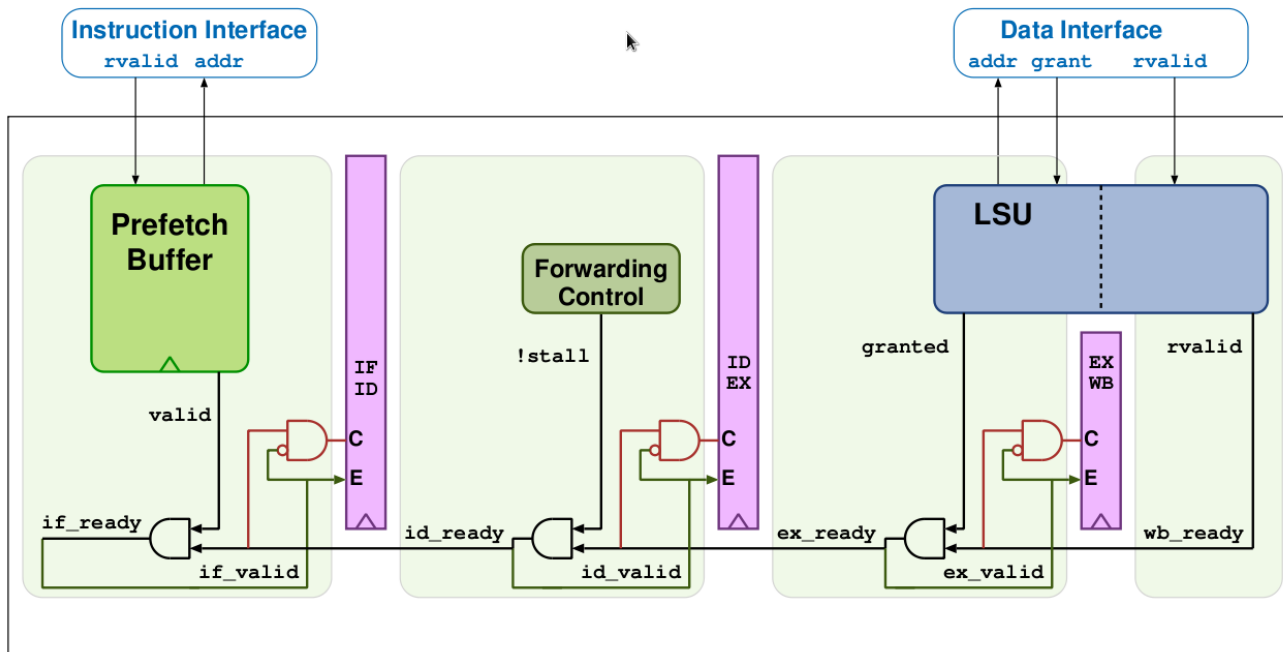


- Instruction memory is word-aligned
 - Does not accept misaligned accesses
- Cross-word instruction needs to be assembled from two words
- If lower half word is compressed, no need to fetch next word already
- Solution: Prefetcher with storage for >2 words
 - We choose 4 words (1 cacheline) for optimal performance and to deal with cache misses



RI5CY: Fully independent pipeline

- Ready & valid signals running left and right respectively
- Pipeline stages can be empty
 - Not possible in OR10N
- Easy to integrate multi-cycle instructions in each stage
 - Similar for “limited” out-of-order execution



RI5CY: Simulation Checker

- Instruction Set Simulator running in parallel using instruction and data access inputs from RTL
- After every cycle write-back data to RF is compared
- ISS serves as golden model
- Could also execute a random instruction stream, since check is automatically and implicitly done
- The simchecker is not activated by default
- To activate:
 - Uncomment ``define simchecker` in riscv_defines.sv
 - Uncomment sv_lib loading in vsim setup files
 - Build and copy rtl_checker from SDK to vsim work directory

RI5CY: Random Stall Injection

- Only on PULPino for the moment
- Random latency (non-synthesizable) introduced to instruction and data accesses
- Checks core interfaces in a way that we would seldom see on the platform
- To activate
 - Uncomment ``define DATA_STALL_RANDOM` and ``define INSTR_STALL_RANDOM` in `config.sv` of PULPino

RI5CY: Tracer

- Similar to objdump output
- Displays also read and written registers
 - x??:read value
 - x??=written value
 - PA: physical address for memory accesses
- Virtual platform follows same output format
- Files automatically generated for RTL simulation
 - trace_core_xx_xx.log


Timestamp	Cyc.	PC	Instr.	Assembler	String	Registers
24530000	ps	1197	0000012c	ffaddce3	bge x27, x26, -8	x27: 001002d4 x26: 001002d4
24510000	ps	1196	00000128	004d0d13	addi x26, x26, 4	x26=001002d4 x26: 001002d0
24590000	ps	1200	00000124	000d2023	sw x0, 0(x26)	x26: 001002d4 PA: 001002d4
24630000	ps	1202	0000012c	ffaddce3	bge x27, x26, -8	x27: 001002d4 x26: 001002d8
24610000	ps	1201	00000128	004d0d13	addi x26, x26, 4	x26=001002d8 x26: 001002d4
24650000	ps	1203	00000130	00000513	addi x10, x0, 0	x10=00000000
24670000	ps	1204	00000134	00100593	addi x11, x0, 1	x11=00000001

RI5CY: Performance Counters

- Events to be counted:
 - #cycles
 - #instructions
 - #ld_stall: load data hazards
 - #jr_stall: number of jump register data hazards
 - #imiss: cycles waiting for instructions
 - #ld
 - #st
 - #jump
 - #branch: total number of branches (w/o jumps)
 - #btaken: branches that were taken
 - #rvc: number of rvc insns
 - #ld_ext: LD to non-tcdm
 - misaligned access counted twice
 - #st_ext: ST to non-tcdm
 - #ld_ext_cyc
 - #st_ext_cyc
 - #tcdm_cont: cycles wasted due to waiting for grants in L1

RI5CY: Performance Counters

- On ASIC only one counter + one register
- On FPGA/RTL sim
 - one counter + one register per metric
- Binary tracer for performance counters is not yet available (for KCG)



Other RISC-V Cores

Z-Scale / V-Scale

- From UC Berkeley
- Their take on a small core
- Z-Scale: Written in Chisel
- V-Scale: Written in Verilog
- Z-Scale & V-Scale virtually identical

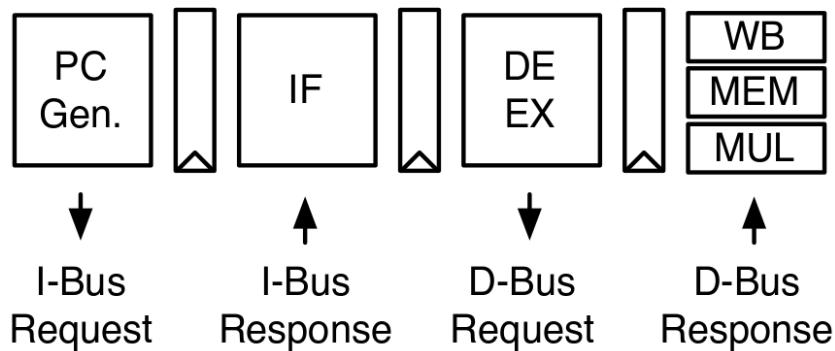
Mini project:

Compare V-Scale/Z-Scale to our core [1]

Semester project:

Design a mini core with lower power consumption than V-Scale, Z-Scale

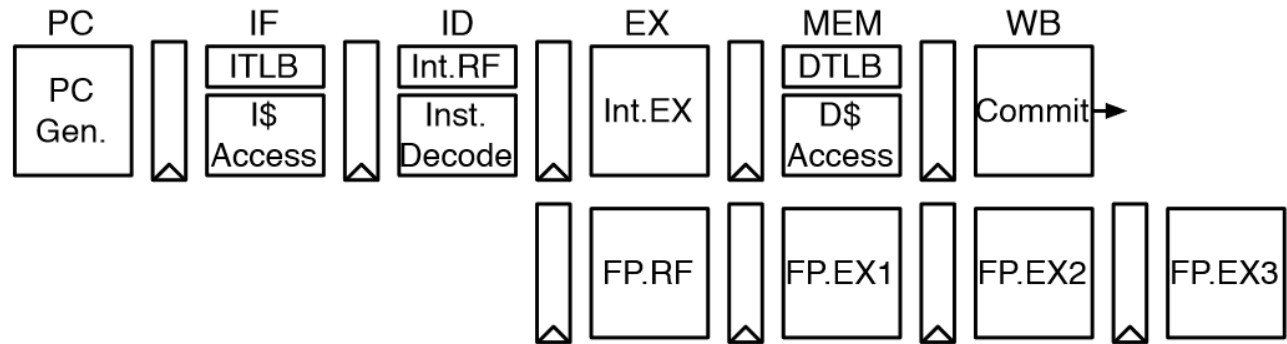
Z-scale Pipeline



Category	ARM Cortex-M0	RISC-V Zscale
ISA	32-bit ARM v6	32-bit RISC-V (RV32IM)
Architecture	Single-Issue In-Order 3-stage	Single-Issue In-Order 3-stage
Performance	0.87 DMIPS/MHz	1.35 DMIPS/MHz
Process	TSMC 40LP	TSMC 40GPLUS
Area w/o Caches	0.0070 mm ²	0.0098 mm ²
Area Efficiency	124 DMIPS/MHz/mm ²	138 DMIPS/MHz/mm ²
Frequency	≤50 MHz	~500 MHz
Voltage (RTV)	1.1 V	0.99 V
Dynamic Power	5.1 μW/MHz	1.8 μW/MHz

Other RISC-V Cores: Rocket

- From UC Berkeley, written in Chisel
- 64-Bit Implementation



Category	ARM Cortex-A5	RISC-V Rocket
ISA	32-bit ARM v7	64-bit RISC-V v2
Architecture	Single-Issue In-Order	Single-Issue In-Order 5-stage
Performance	1.57 DMIPS/MHz	1.72 DMIPS/MHz
Process	TSMC 40GPLUS	TSMC 40GPLUS
Area w/o Caches	0.27 mm ²	0.14 mm ²
Area with 16K Caches	0.53 mm ²	0.39 mm ²
Area Efficiency	2.96 DMIPS/MHz/mm ²	4.41 DMIPS/MHz/mm ²
Frequency	>1GHz	>1GHz
Dynamic Power	<0.08 mW/MHz	0.034 mW/MHz

Other RISC-V Cores

BOOM: Berkeley Out-of-Order Processor

- From UC Berkeley, written in Chisel
- Parametrizable for Dual-Issue/Quad-Issue

Category	ARM Cortex-A9	RISC-V BOOM-2w
ISA	32-bit ARM v7	64-bit RISC-V v2 (RV64G)
Architecture	2 wide, 3+1 issue Out-of-Order 8-stage	2 wide, 3 issue Out-of-Order 6-stage
Performance	3.59 CoreMarks/MHz	3.91 CoreMarks/MHz
Process	TSMC 40GPLUS	TSMC 40GPLUS
Area with 32K caches	~2.5 mm ²	~1.00 mm ²
Area efficiency	1.4 CoreMarks/MHz/mm ²	3.9 CoreMarks/MHz/mm ²
Frequency	1.4 GHz	1.5 GHz
Power	0.5-1.9 W (2 cores + L2) @ TSMC 40nm, 0.8-2.0 GHz	0.25 W (1 core + L1) @ TSMC 45nm, 1 GHz

+9%!

Master project: [1]
Design and implement a VLIW-architecture supporting the RISC-V ISA

Mini Project:
Initial design consideration for the implementation (pro/cons evaluation)

note:
not to scale



Exercise session:

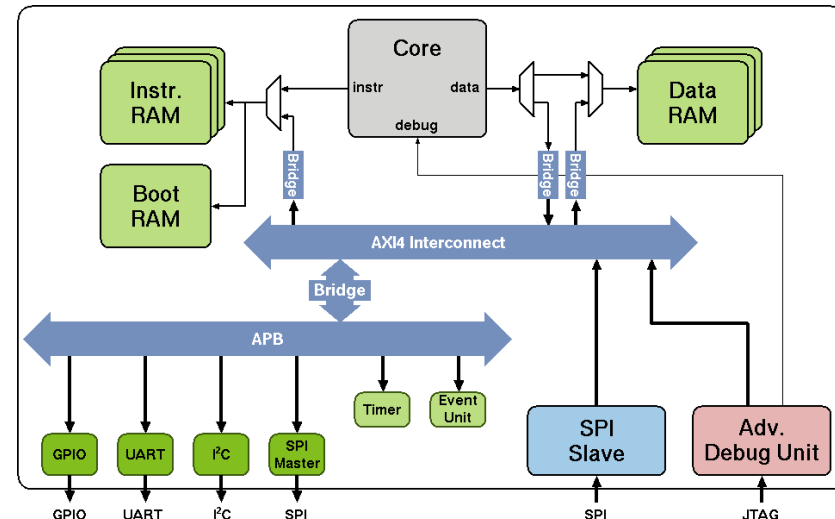
- In the exercise we are going to cover:

- How to compile and run an application using:
 - The open source Pulpino platform

Pulpino architecture

- Impact of the new instructions:

- Hardware loops
- Pre/post increment
- Vector support
- Dot product
- Shuffle instruction



- How to use highly optimized kernels in programs
 - Convolutions
- Comparison of RISC-V to ARM Cortex M4

Q&A

