

MULTICORE PROGRAMMING WITH OPENMP

Ing. Andrea Marongiu


a.marongiu@iis.ee.ethz.ch

Programming model: OpenMP


- De-facto standard for the **shared memory** programming model
- A collection of **compiler directives**, **library routines** and **environment variables**
- Easy to specify parallel execution within a **serial code**
- Requires **special support** in the compiler
- Generates calls to **threading libraries** (e.g. pthreads)
- Focus on **loop-level** parallel execution
- Popular in high-end embedded

Fork/Join Parallelism

Sequential program



Parallel program



- Initially only master thread is active
- Master thread executes sequential code
- Fork: Master thread creates or awakens additional threads to execute parallel code
- Join: At the end of parallel code created threads are suspended upon **barrier** synchronization

Pragmas

- **Pragma**: a compiler directive in C or C++
- Stands for “pragmatic information”
- A way for the programmer to communicate with the compiler
- Compiler free to ignore pragmas: original sequential semantic is not altered
- Syntax:

#pragma omp *<rest of pragma>*

Components of OpenMP

Directives

- ❖ Parallel regions
 - `#pragma omp parallel`
- ❖ Work sharing
 - `#pragma omp for`
 - `#pragma omp sections`
- ❖ Synchronization
 - `#pragma omp barrier`
 - `#pragma omp critical`
 - `#pragma omp atomic`

Runtime Library

Clauses

- ❖ Data scope attributes
 - `private`
 - `shared`
 - `reduction`
- ❖ Loop scheduling
 - `static`
 - `dynamic`

- ❖ Thread Forking/Joining
 - `omp_parallel_start()`
 - `omp_parallel_end()`
- ❖ Loop scheduling
- ❖ Thread IDs
 - `omp_get_thread_num()`
 - `omp_get_num_threads()`

Outlining parallelism

The `parallel` directive

- Fundamental construct to outline parallel computation within a sequential program
- Code within its scope is **replicated** among threads
- Defers implementation of parallel execution to the runtime (machine-specific, e.g. **`pthread_create`**)

**A sequential program..
..is easily parallelized**

```
int main()  
{  
  #pragma omp parallel  
  {  
    printf ("\nHello world!");  
  }  
}
```

```
int main()  
{  
  omp_parallel_start(&parfun, ...);  
  parfun();  
  omp_parallel_end();  
}  
  
int parfun(...)  
{  
  printf ("\nHello world!");  
}
```

#pragma omp parallel

Code originally contained within the scope of the pragma is outlined to a new function within the compiler

```
int main()  
{  
#pragma omp parallel  
{  
    printf ("\nHello world!");  
}  
}
```

```
int main()  
{  
    omp_parallel_start(&parfun, ...);  
    parfun();  
    omp_parallel_end();  
}  
  
int parfun(...)  
{  
    printf ("\nHello world!");  
}
```

#pragma omp parallel

The `#pragma` construct in the **main** function is replaced with function calls to the runtime library

```
int main()  
{  
#pragma omp parallel  
{  
    printf ("\nHello world!");  
}  
}
```

```
int main()  
{  
    omp_parallel_start(&parfun, ...);  
    parfun();  
    omp_parallel_end();  
}  
  
int parfun(...)  
{  
    printf ("\nHello world!");  
}
```


#pragma omp parallel


First we call the runtime to fork new threads, and pass them a pointer to the function to execute in parallel

```
int main()  
{  
    #pragma omp parallel  
    {  
        printf ("\nHello world!");  
    }  
}
```

```
int main()  
{  
    omp_parallel_start(&parfun, ...);  
    parfun();  
    omp_parallel_end();  
}  
  
int parfun(...)  
{  
    printf ("\nHello world!");  
}
```

#pragma omp parallel

Then the master itself calls
the parallel function




```
int main()  
{  
    #pragma omp parallel  
    {  
        printf ("\nHello world!");  
    }  
}
```

```
int main()  
{  
    omp_parallel_start(&parfun, ...);  
    parfun();  
    omp_parallel_end();  
}  
  
int parfun(...)  
{  
    printf ("\nHello world!");  
}
```

#pragma omp parallel

Finally we call the runtime to synchronize threads with a barrier and suspend them



```
int main()  
{  
    #pragma omp parallel  
    {  
        printf ("\nHello world!");  
    }  
}
```

```
int main()  
{  
    omp_parallel_start(&parfun, ...);  
    parfun();  
    omp_parallel_end();  
}  
  
int parfun(...)  
{  
    printf ("\nHello world!");  
}
```

#pragma omp parallel

Data scope attributes

```
int main()  
{  
    int id;  
    int a = 5;  
    #pragma omp parallel  
    {  
        id = omp_get_thread_num();  
        if (id == 0)  
            printf ("Master: a = %d.", a*2);  
        else  
            printf ("Slave: a = %d.", a);  
    }  
}
```

Call runtime to get thread ID:
Every thread sees a different value

Master and slave threads
access the same variable **a**

A slightly more complex example

#pragma omp parallel

Data scope attributes

```
int main()
```

```
{
```

```
    int id;
```

```
    int a = 5;
```

```
    #pragma omp parallel
```

```
    {
```

```
        id = omp_get_thread_num();
```

```
        if (id == 0)
```

```
            printf ("Master: a = %d.", a);
```

```
        else
```

```
            printf ("Slave: a = %d.", a);
```

```
    }
```

```
}
```

Call runtime to get thread ID:

Every thread sees a different value

How to inform the compiler
about these different
behaviors?

Master and slave threads
access same variable **a**

A slightly more complex example

#pragma omp parallel

Data scope attributes

```
int main()  
{  
    int id;  
    int a = 5;  
    #pragma omp parallel shared (a) private (id)  
    {  
        id = omp_get_thread_num();  
        if (id == 0)  
            printf ("Master: a = %d.", a*2);  
        else  
            printf ("Slave: a = %d.", a);  
    }  
}
```

Insert code to retrieve the address
of the shared object from within
each parallel thread

Allow symbol privatization:
Each thread contains a
private copy of this variable

A slightly more complex example

#pragma omp parallel

Data scope attributes

```
int main()  
{  
    int id;  
    int a =  
#pragma omp  
{  
    id = omp  
    if (id =  
        printf  
    else  
        printf  
    }  
}
```

Correctness issues

What if:

- **a** was not marked for shared access?
- **id** was not marked for private access?

ert code to retrieve the address
e shared object from within
parallel thread

(id)

low symbol privatization:
ch thread contains a
rate copy of this variable

A slightly more complex example

More data sharing clauses

- firstprivate
 - ▣ copyin, private storage
- lastprivate
 - ▣ Copyout, private storage

Sharing work among threads

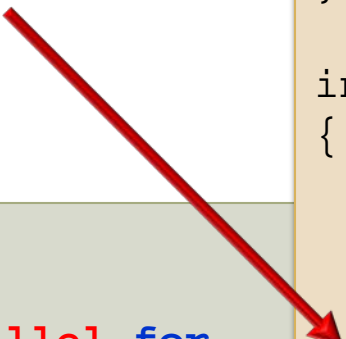
The **for** directive

- The **parallel** pragma instructs every thread to execute all of the code inside the block
- If we encounter a **for** loop that we want to divide among threads, we use the **for** pragma

```
#pragma omp for
```

#pragma omp for

The code of the **for** loop is moved inside the outlined function.

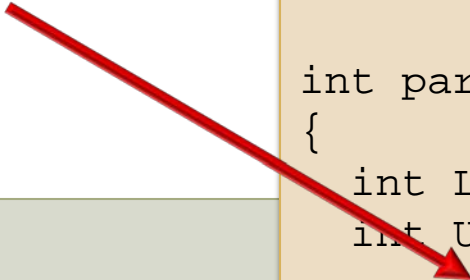


```
int main()  
{  
  #pragma omp parallel for  
  {  
    for (i=0; i<10; i++)  
      a[i] = i;  
  }  
}
```

```
int main()  
{  
  omp_parallel_start(&parfun, ...);  
  parfun();  
  omp_parallel_end();  
}  
  
int parfun(...)  
{  
  int LB = ...;  
  int UB = ...;  
  for (i=LB; i<UB; i++)  
    a[i] = i;  
}
```

#pragma omp for

Every thread works on a different subset of the iteration space..

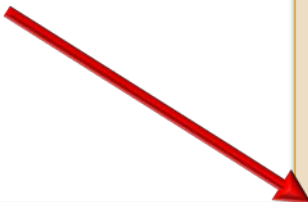


```
int main()  
{  
  #pragma omp parallel for  
  {  
    for (i=0; i<10; i++)  
      a[i] = i;  
  }  
}
```

```
int main()  
{  
  omp_parallel_start(&parfun, ...);  
  parfun();  
  omp_parallel_end();  
}  
  
int parfun(...)  
{  
  int LB = ...;  
  int UB = ...;  
  
  for (i=LB; i<UB; i++)  
    a[i] = i;  
}
```

#pragma omp for

..since lower and upper boundaries (LB, UB) are computed locally



```
int main()  
{  
  #pragma omp parallel for  
  {  
    for (i=0; i<10; i++)  
      a[i] = i;  
  }  
}
```

```
int main()  
{  
  omp_parallel_start(&parfun, ...);  
  parfun();  
  omp_parallel_end();  
}  
  
int parfun(...)  
{  
  int LB = ...;  
  int UB = ...;  
  
  for (i=LB; i<UB; i++)  
    a[i] = i;  
}
```

The schedule clause

Static Loop Partitioning


```
#pragma omp for schedule(static)
{
    for (i=0; i<12; i++)
        a[i] = i;
}
```

Es. 12 iterations (N), 4 threads (Nthr)

DATA CHUNK

$$C = \text{ceil} \left(\frac{N}{N_{\text{thr}}} \right)$$

**3
iterations
thread**



Thread ID (TID)

0

1

2

3

LOWER BOUND

$$LB = C * TID$$

0

3

6

9

UPPER BOUND

$$UB = \min \{ [C * (TID + 1)], N \}$$

3

6

9

12

The schedule clause

Static Loop Partitioning

Es. 12 iterations (N), 4 threads (Nthr)

```
#pragma omp for schedule(static)
{
    for (i=0; i<12; i++)
        a[i] = i;
}
```

Useful for:

- Simple, regular loops
- Iterations with equal duration

DATA CHUNK

$$\text{chunk} \left(\frac{N}{N_{\text{thr}}} \right)$$

**3
iterations
thread**

**What happens with
static scheduling
when iterations have
different duration?**

LOWER BOUND

UPPER BOUND

$\text{chunk} \cdot (\text{TID} + 1), N\}$

	0	1	2	3
	0	3	6	9
	3	6	9	12

The schedule clause

Static Loop Partitioning


```
#pragma omp for schedule(static)
{
    for (i=0; i<12; i++)
    {
        int start = rand();
        int count = 0;

        while (start++ < 256)
            count++;

        a[count] = foo();
    }
}
```

A variable amount of work
in each iteration

UNBALANCED
workloads



The `schedule` clause

Dynamic Loop Partitioning

```
#pragma omp for schedule(dynamic)
{
    for (i=0; i<12; i++)
    {
        int start = rand();
        int count = 0;

        while (start++ < 256)
            count++;

        a[count] = foo();
    }
}
```



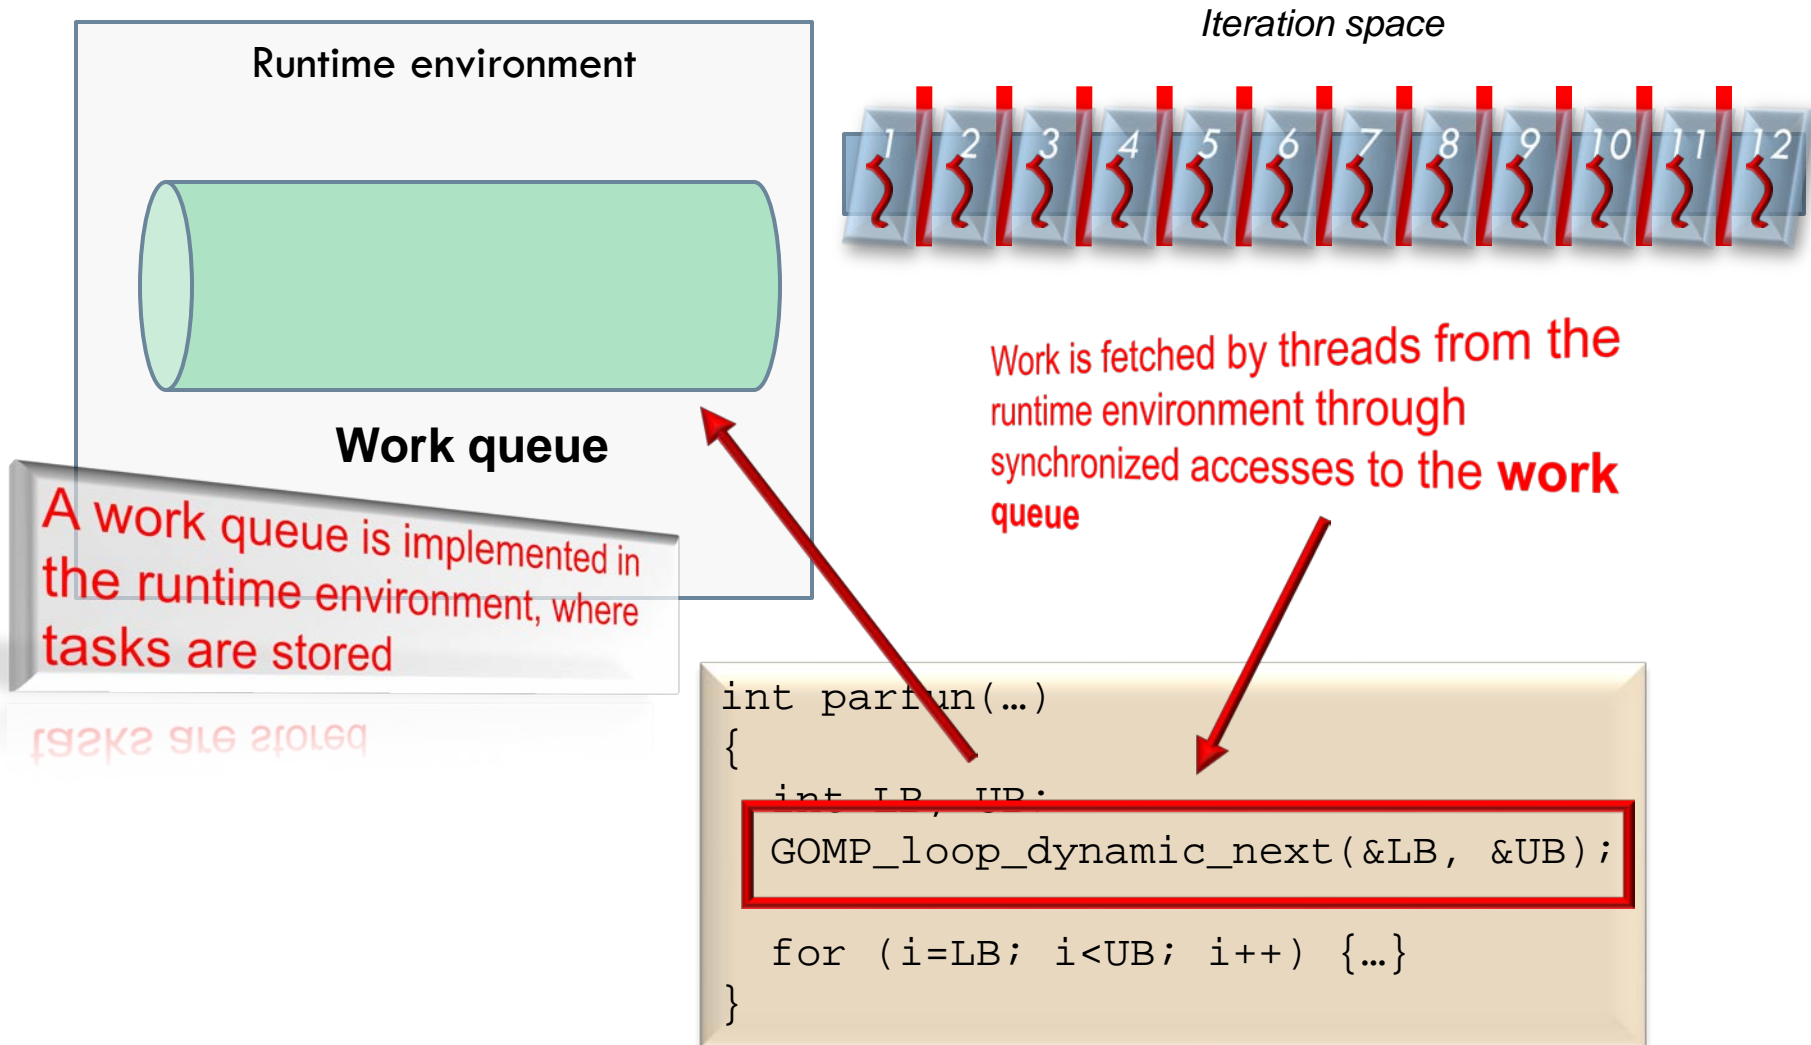
A thread is generated for
every single iteration

Iteration space



The `schedule` clause




Dynamic Loop Partitioning



The schedule clause

Dynamic Loop Partitioning

Remember results with static scheduling..



Speedup!


BALANCED
workloads

Parallelization granularity

Iteration chunking

● Fine-grain Parallelism

- Best opportunities for load balancing, but..
- Small amounts of computational work between parallelism computation stages
- Low computation to parallelization ratio → High parallelization overhead



● Coarse-grain Parallelism

- Harder to load balance efficiently, but..
- Large amounts of computational work between parallelism computation stages
- High computation to parallelization ratio → Low parallelization overhead
- Harder to load balance



The `schedule` clause

Dynamic Loop Partitioning

Iteration space

```
#pragma omp for schedule(dynamic, 1)
{
    for (i=0; i<12; i++)
    {
        int start = rand();
        int count = 0;



        while (start++ < 256)
            count++;

        a[count] = foo();
    }
}
```



The `schedule` clause

Dynamic Loop Partitioning



Runtime scheduling primitive is invoked
at every iteration


If granularity of WORK is very
small the overhead for fine
chunking is significant

```
int parfun(...)
{
    int LB, UB;
    GOMP_loop_dynamic_next(&LB, &UB);
    for (i=LB; i<UB; i++) {WORK}
}
```



The schedule clause

Dynamic Loop Partitioning

**Remember results with
static scheduling..**



Iteration space





The schedule clause

Dynamic Loop Partitioning

chunking overhead

Iteration space



The `schedule` clause

Dynamic Loop Partitioning

Iteration space

```
#pragma omp for schedule(dynamic, 2)
{
    for (i=0; i<12; i++)
    {
        int start = rand();
        int count = 0;



        while (start++ < 256)
            count++;

        a[count] = foo();
    }
}
```



The `schedule` clause

Dynamic Loop Partitioning



Runtime scheduling primitive is invoked every two iterations (half the times of previous case)

WORK is repeated twice among two scheduling events. Overhead gets amortized



```
int parfun(...)
{
    int LB, UB;
    GOMP_loop_dynamic_next(&LB, &UB);
    for (i=LB; i<UB; i++) {WORK}
}
```

The schedule clause

Dynamic Loop Partitioning

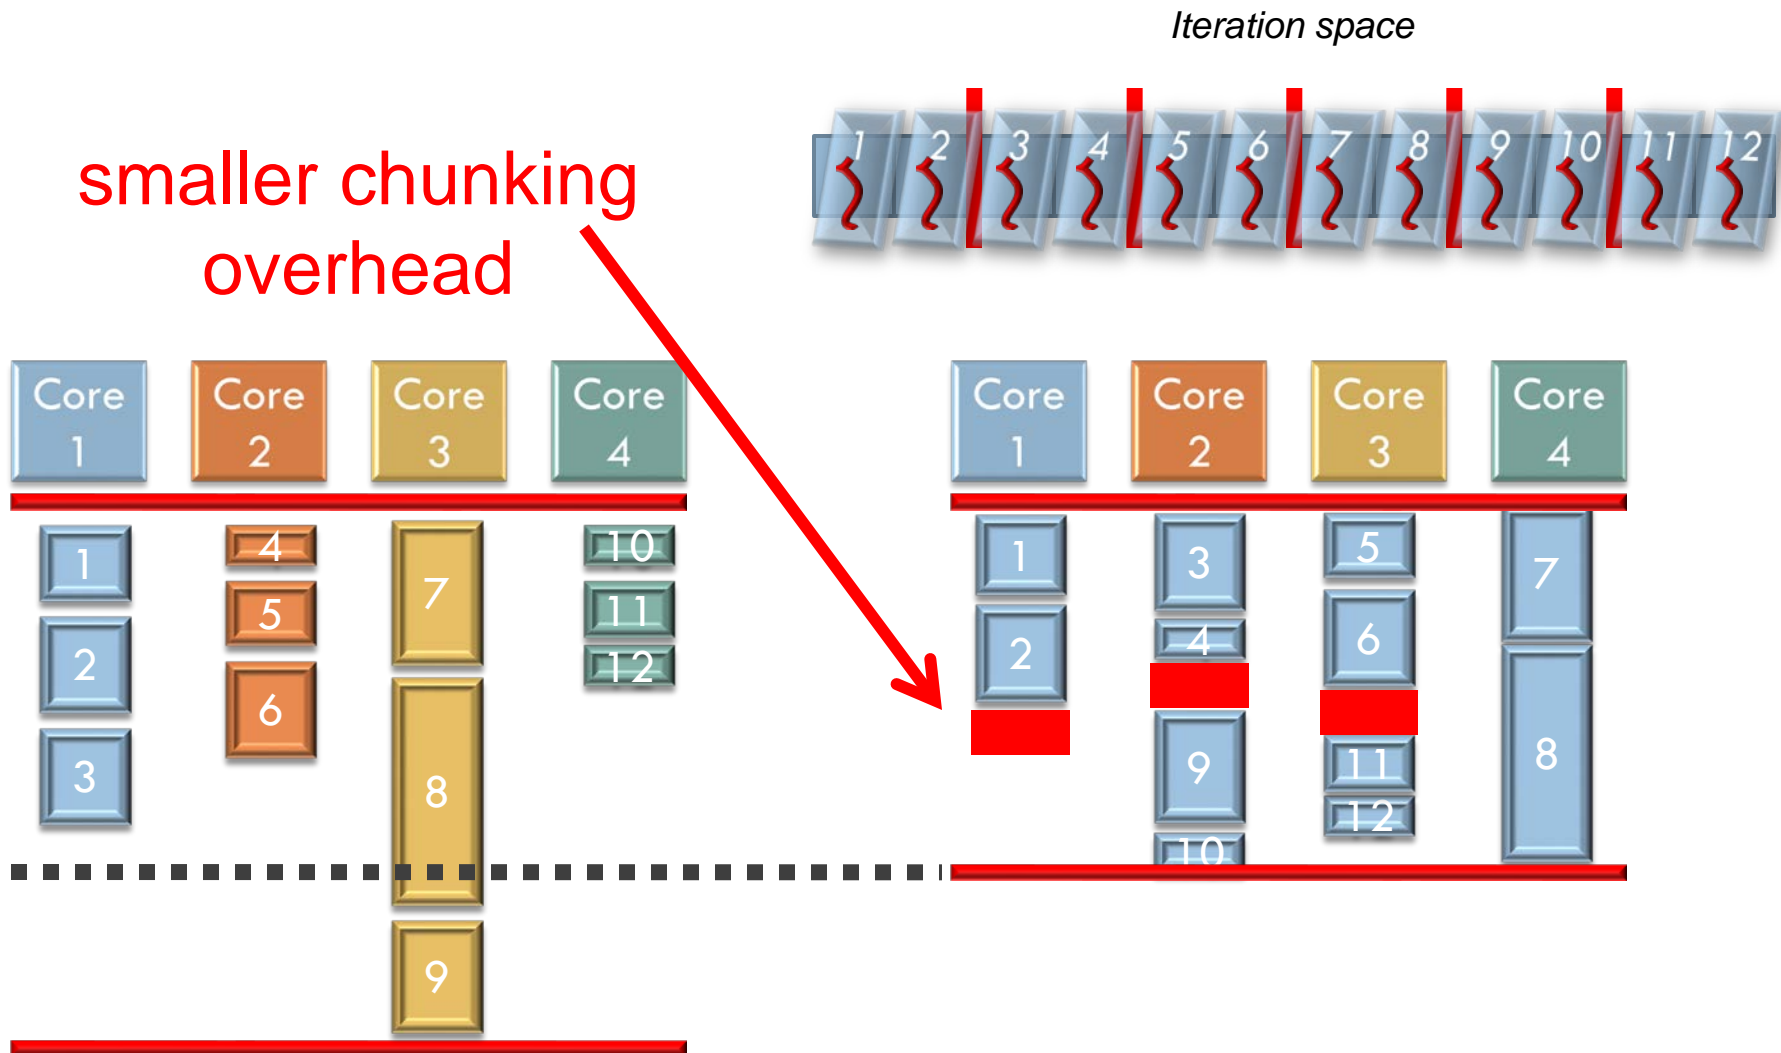
**Remember results with
static scheduling..**

Iteration space



The schedule clause

Dynamic Loop Partitioning



More scheduling clauses

□ **schedule (guided[, *chunk*])**

- ▣ Threads dynamically grab blocks of iterations. The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds.

□ **schedule (runtime[, *chunk*])**

- ▣ Schedule and chunk size taken from the OMP_SCHEDULE environment variable (or the runtime library ... for OpenMP 3.0)

Sharing work among threads

The `sections` directive

- The `for` pragma allows to exploit **data parallelism** in loops
- OpenMP also provides directives to exploit **task parallelism**

```
#pragma omp sections
```

SPMD VS MPMD

Recall..


- SPMD (single program, multiple data)
 - ▣ Processors execute the same stream of instructions over different data
 - ▣ `#pragma omp for`
- MPMD (multiple program, multiple data)
 - ▣ Processors execute different streams of instructions over (possibly) different data
 - ▣ `#pragma omp sections`
 - ▣ `#pragma omp task`

Task Parallelism Example

```
int main()  
{  
  
    v = alpha();  
    w = beta ();  
  
    y = delta ();  
    x = gamma (v, w);  
    z = epsilon (x, y));  
  
    printf ("%f\n", z);  
}
```

**FIRST
SOLUTION**


Identify independent nodes
in the task graph, and outline
parallel computation with the
sections directive



Task Parallelism Example

```
int main()  
{  
#pragma omp parallel sections {  
    v = alpha();  
    w = beta ();  
}  
#pragma omp parallel sections {  
    y = delta ();  
    x = gamma (v, w);  
}  
    z = epsilon (x, y));  
  
    printf ("%f\n", z);  
}
```


Barriers implied here!



Task Parallelism Example

```
int main()  
{  
    #pragma omp parallel sections {  
        v = alpha();  
        w = beta ();  
    }  
    #pragma omp parallel sections {  
        y = delta ();  
        x = gamma (v, w);  
    }  
    z = epsilon (x, y);  
    printf ("%f\n", z);  
}
```


Each parallel task within a
sections block identifies a
section



Task Parallelism Example

```
int main()  
{  
#pragma omp parallel sections {  
    #pragma omp section  
        v = alpha();  
    #pragma omp section  
        w = beta ();  
}  
#pragma omp parallel sections {  
    #pragma omp section  
        y = delta ();  
    #pragma omp section  
        x = gamma (v, w);  
}  
    z = epsilon (x, y));  
  
    printf ("%f\n", z);  
}
```

Each parallel task within a
sections block identifies a
section




Task Parallelism Example

```
int main()  
{  
  
    v = alpha();  
  
    w = beta ();  
  
    y = delta ();  
  
  
    x = gamma (v, w);  
  
    z = epsilon (x, y));  
  
    printf ("%f\n", z);  
}
```

**SECOND
SOLUTION**


Identify independent nodes
in the task graph, and outline
parallel computation with the
sections directive



Task Parallelism Example

```
int main()  
{  
    #pragma omp parallel sections {  
  
        v = alpha();  
  
        w = beta ();  
  
        y = delta ();  
    }  
  
    x = gamma (v, w);  
  
    z = epsilon (x, y));  
  
    printf ("%f\n", z);  
}
```


Barrier implied here!



Task Parallelism Example

```
int main()  
{  
    #pragma omp parallel sections {  
  
        v = alpha();  
  
        w = beta ();  
  
        y = delta ();  
    }  
  
    x = gamma (v, w);  
  
    z = epsilon (x, y));  
  
    printf ("%f\n", z);  
}
```


Each parallel task within a
sections block identifies a
section



Task Parallelism Example

```
int main()  
{  
    #pragma omp parallel sections {  
        #pragma omp section  
        v = alpha();  
        #pragma omp section  
        w = beta ();  
        #pragma omp section  
        y = delta ();  
    }  
  
    x = gamma (v, w);  
  
    z = epsilon (x, y));  
  
    printf ("%f\n", z);  
}
```

Each parallel task within a
sections block identifies a
section



#pragma omp barrier

- Most important **synchronization** mechanism in shared memory fork/join parallel programming
- All threads participating in a parallel region wait until everybody has finished before computation flows on
- This prevents later stages of the program to work with **inconsistent** shared data
- It is implied at the end of **parallel** constructs, as well as **for** and **sections** (unless a **nowait** clause is specified)


#pragma omp critical

- **Critical Section:** a portion of code that only one thread at a time may execute
- We denote a critical section by putting the pragma

#pragma omp critical

in front of a block of C code

π -finding code example




```
double area, pi, x;
int i, n;
#pragma omp parallel for private(x) \
    shared(area)
{
    for (i=0; i<n; i++) {
        x = (i + 0.5)/n;
        area += 4.0/(1.0 + x*x);
    }
    pi = area/n;
}
```

Synchronize accesses to
shared variable **area** to
avoid inconsistent results

Race condition


- Ensure atomic updates of the shared variable **area** to avoid a **race condition** in which one process may “race ahead” of another and ignore changes



Race condition (Cont'd)


- Thread A reads “11.667” into a local register
- Thread B reads “11.667” into a local register
- Thread A updates area with “11.667+3.765”
- Thread B ignores write from thread A and updates area with “11.667 + 3.563”

time



π -finding code example

```
double area, pi, x;
int i, n;
#pragma omp parallel for private(x) shared(area)
{
    for (i=0; i<n; i++) {
        x = (i + 0.5)/n;
        #pragma omp critical
        area += 4.0/(1.0 + x*x);
    }
}
pi = area/n;
```



#pragma omp critical protects the code within its scope by acquiring a lock before entering the critical section and releasing it after execution

Correctness, not performance!

- As a matter of fact, using locks makes execution **sequential**
- To dim this effect we should try use **fine grained locking** (i.e. make critical sections as small as possible)
- A simple instruction to compute the value of area in the previous example is translated into many more simpler instructions within the compiler!
- The programmer is not aware of the real **granularity** of the critical section

formance!

```
suif_start.omp_fn.0 (.omp_data_i)
{
  <bb 2>:
    D.2605 = .omp_data_i->n;
    D.2615 = __builtin_omp_get_num_threads ();
    D.2616 = __builtin_omp_get_thread_num ();

    ...
    D.2623 = MIN_EXPR <D.2622, D.2605>;
    if (D.2621 >= D.2623) goto <L4>; else goto <L2>;

  <L4>::
    return;

  <L2>::
    D.2586 = (double) i;
    D.2587 = D.2586 + 5.0e-1;
    D.2605 = .omp_data_i->n;
    D.2606 = D.2605;
    D.2588 = (double) D.2606;
    x = D.2587 / D.2588;
    __builtin_GOMP_critical_start ();
    D.2589 = x * x;
    D.2590 = D.2589 + 1.0e+0;
    D.2591 = 4.0e+0 / D.2590;
    D.2607 = .omp_data_i->area;
    D.2608 = D.2607;
    D.2609 = D.2591 + D.2608;
    .omp_data_i->area = D.2609;
    __builtin_GOMP_critical_end ();
    i = i + 1;
    D.2627 = i < D.2626;
    if (D.2627) goto <L2>; else goto <L4>;
}
```

Execution sequential

This is a dump of the intermediate representation of the program within the compiler

```

suif_start.omp_fn.0 (.omp_data_i)
{
  <bb 2>:
    D.2605 = .omp_data_i->n;
    D.2615 = __builtin_omp_get_num_threads ();
    D.2616 = __builtin_omp_get_thread_num ();

    ...
    D.2623 = MIN_EXPR <D.2622, D.2605>;
    if (D.2621 >= D.2623) goto <L4>; else goto <L2>;

  <L4>::
    return;

  <L2>::
    D.2586 = (double) i;
    D.2587 = D.2586 + 5.0e-1;
    D.2605 = .omp_data_i->n;
    D.2606 = D.2605;
    D.2588 = (double) D.2606;
    x = D.2587 / D.2588;
    __builtin_GOMP_critical_start ();
    D.2589 = x * x;
    D.2590 = D.2589 + 1.0e+0;
    D.2591 = 4.0e+0 / D.2590;
    D.2607 = .omp_data_i->area;
    D.2608 = D.2607;
    D.2609 = D.2591 + D.2608;
    .omp_data_i->area = D.2609;
    __builtin_GOMP_critical_end ();
    i = i + 1;
    D.2627 = i < D.2626;
    if (D.2627) goto <L2>; else goto <L4>;
}

```

Performance!

Execution sequential
 fine-grained locking (i.e.
 e)

of area in the
 / more simpler

granularity of the

This is how the compiler
 represents the instruction

$area += 4.0 / (1.0 + x * x);$

formance!

**THIS IS DONE AT EVERY
LOOP ITERATION!**

```
_suif_start.omp_fn.0 (.omp_data_i)
```

```
<bb 2>:
```

```
D.2605 = .omp_data_i->n;
```

```
D.2615 = __builtin_omp_get_num_threads();
```

```
D.2616 = __builtin_omp_get_thread_num();
```

```
***  
D.2623 = MIN_EXPR <D.2622, D.2605>;
```

```
if (D.2621 >= D.2623) goto <L4>; else goto <L2>;
```

```
<L4>;
```

```
return;
```

```
<L2>;
```

```
D.2586 = (double) i;
```

```
D.2587 = D.2586 + 5.0e-1;
```

```
D.2605 = .omp_data_i->n;
```

```
D.2606 = D.2605;
```

```
D.2588 = (double) D.2606;
```

```
x = D.2587 / D.2588;
```

```
__builtin_GOMP_critical_start();
```

```
D.2589 = x * x;
```

```
D.2590 = D.2589 + 1.0e+0;
```

```
D.2591 = 4.0e+0 / D.2590;
```

```
D.2607 = .omp_data_i->area;
```

```
D.2608 = D.2607;
```

```
D.2609 = D.2591 + D.2608;
```

```
__omp_data_i->area = D.2609;
```

```
__builtin_GOMP_critical_end();
```

```
i = i + 1;
```

```
D.2627 = i < D.2626;
```

```
if (D.2627) goto <L2>; else goto <L4>;
```

call runtime to acquire lock

Lock-protected
operations
(*critical section*)

call runtime to release lock

Execution sequential
grained locking (i.e.
e)

of area in the
more simpler

granularity of the


3,0-1

π -finding code example

```
double area, pi, x;
int i, n;
#pragma omp parallel for \
    private(x) \
    shared(area)
{
    for (i=0; i<n; i++) {
        x = (i + 0.5)/n;
        #pragma omp critical
        area += 4.0/(1.0 + x*x);
    }
    pi = area/n;
}
```



Waiting for lock



Correctness, not performance!

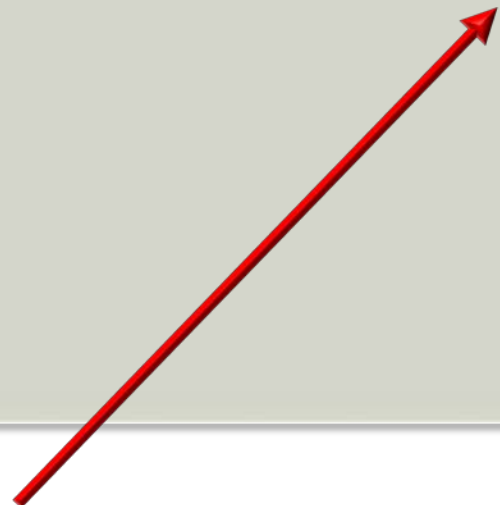
- A programming pattern such as `area += 4.0/(1.0 + x*x);` in which we:
 - ▣ *Fetch the value of an operand*
 - ▣ *Add a value to it*
 - ▣ *Store the updated value*is called a **reduction**, and is commonly supported by parallel programming APIs

- OpenMP takes care of storing partial results in **private variables** and combining partial results after the loop

Correctness, not performance!

```
double area, pi, x;
int i, n;
#pragma omp parallel for private(x) shared(area) reduction(+:area)
{
    for (i=0; i<n; i++) {
        x = (i + 0.5)/n;

        area += 4.0/(1.0 + x*x);
    }
}
pi = area/n;
```



The **reduction** clause instructs the compiler to create **private** copies of the **area** variable for every thread. At the end of the loop partial sums are combined on the shared **area** variable

Performance!

```
suif_start.omp_fn.0 (.omp_data_i)
{
  <bb 2>:
    area = 0.0;
    D.2605 = .omp_data_i->n;
    D.2615 = __builtin_omp_get_num_threads ();
    D.2616 = __builtin_omp_get_thread_num ();

    ...
    D.2623 = MIN_EXPR <D.2622, D.2605>;
    if (D.2621 >= D.2623) goto <L4>; else goto <L2>;

```

```
<L4>::
  builtin_GOMP_atomic_start ();
  D.2607 = &.omp_data_i->area;
  D.2608 = *D.2607;
  D.2609 = D.2608 + area;
  *D.2607 = D.2609;
  __builtin_GOMP_atomic_end ();
  return;

```

```
<L2>::
  D.2586 = (double) i;
  D.2587 = D.2586 + 5.0e-1;
  D.2605 = .omp_data_i->n;
  D.2606 = D.2605;
  D.2588 = (double) D.2606;
  D.2587 / D.2588;
  x = D.2587 / D.2588;
  D.2589 = x * x;
  D.2590 = D.2589 + 1.0e+0;
  D.2591 = 4.0e+0 / D.2590;
  D.2591 + area;
  area = D.2591 + area;
  i = i + 1;
  D.2627 = i < D.2626;
  if (D.2627) goto <L2>; else goto <L4>;

```

Shared variable is only updated at the end of the loop, when partial sums are computed

shared(area) reduction(+:area)

Each thread computes partial sums on private copies of the reduction variable

The r
the a
comb

er to create **private** copies of
nd of the loop partial sums are

Performance!

double

int

#pragma

{

for

a

}

pi

```
_suif_start.omp_fn.0 (.omp_data_i)
{
  <bb 2>:
    area = 0.0;
    D.2605 = .omp_data_i->n;
    D.2615 = __builtin_omp_get_num_threads ();
    D.2616 = __builtin_omp_get_thread_num ();

    ...
    D.2623 = MIN_EXPR <D.2622, D.2605>;
    if (D.2621 >= D.2623) goto <L4>; else goto <L2>;

```

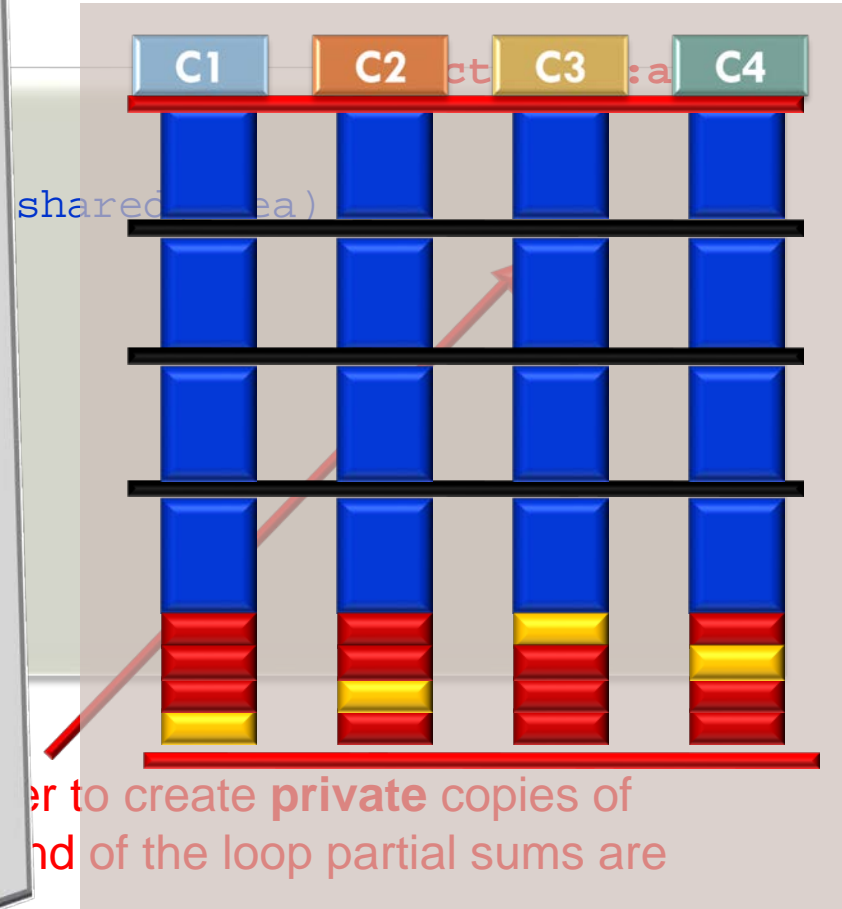
```
<L4>::
  builtin_GOMP_atomic_start ();
  D.2607 = &.omp_data_i->area;
  D.2608 = *D.2607;
  D.2609 = D.2608 + area;
  *D.2607 = D.2609;
  __builtin_GOMP_atomic_end ();
  return;

```

```
<L2>::
  D.2586 = (double) i;
  D.2587 = D.2586 + 5.0e-1;
  D.2605 = .omp_data_i->n;
  D.2606 = D.2605;
  D.2588 = (double) D.2606;
  x = D.2587 / D.2588;
  D.2589 = x * x;
  D.2590 = D.2589 + 1.0e+0;
  D.2591 = 4.0e+0 / D.2590;
  area = D.2591 + area;
  i = i + 1;
  D.2627 = i < D.2626;
  if (D.2627) goto <L2>; else goto <L4>;

```

The r
the a
comb



More worksharing constructs

The **master** directive

- The **master** construct denotes a structured block that is only executed by the master thread.
- The other threads just skip it (no synchronization is implied).

```
#pragma omp parallel
{
    do_many_things();
    #pragma omp master
    { exchange_boundaries(); }

    #pragma omp barrier
    do_many_other_things();
}
```

More worksharing constructs

The **single** directive

- The **single** construct denotes a block of code that is executed by only one thread (not necessarily the master thread).
- A barrier is implied at the end of the single block (can remove the barrier with a **nowait** clause).

```
#pragma omp parallel
{
    do_many_things();
    #pragma omp single
    { exchange_boundaries(); }

    #pragma omp barrier
    do_many_other_things();
}
```

Recap: TASK parallelism in OpenMP 2.5

The `sections` directive

- The `for` pragma allows to exploit `data parallelism` in loops
- OpenMP 2.5 also provides a directive to exploit `task parallelism`


```
#pragma omp sections
```


Task Parallelism Example

```
int main()  
{  
  
    v = alpha();  
    w = beta ();  
  
    y = delta ();  
    x = gamma (v, w);  
    z = epsilon (x, y));  
  
    printf ("%f\n", z);  
}
```

**FIRST
SOLUTION**


Identify independent nodes
in the task graph, and outline
parallel computation with the
sections directive



Task Parallelism Example

```
int main()  
{  
  #pragma omp parallel sections {  
    v = alpha();  
    w = beta ();  
  }  
  #pragma omp parallel sections {  
    y = delta ();  
    x = gamma (v, w);  
  }  
  z = epsilon (x, y);  
  printf ("%f\n", z);  
}
```


Barriers implied here!



Task Parallelism Example

```
int main()  
{  
    #pragma omp parallel sections {  
        v = alpha();  
        w = beta ();  
    }  
    #pragma omp parallel sections {  
        y = delta ();  
        x = gamma (v, w);  
    }  
    z = epsilon (x, y);  
    printf ("%f\n", z);  
}
```


Each parallel task within a
sections block identifies a
section



Task Parallelism Example

```
int main()  
{  
  #pragma omp parallel sections {  
    #pragma omp section  
      v = alpha();  
    #pragma omp section  
      w = beta ();  
  }  
  #pragma omp parallel sections {  
    #pragma omp section  
      y = delta ();  
    #pragma omp section  
      x = gamma (v, w);  
  }  
  z = epsilon (x, y));  
  
  printf ("%f\n", z);  
}
```

Each parallel task within a
sections block identifies a
section



Task parallelism

- The **sections** directive allows a very limited form of task parallelism
- All tasks must be statically outlined in the code
 - ▣ What if a functional loop (**while**) body is identified as a task?
 - **Unrolling?** Not feasible for high iteration count
 - ▣ What if recursion is used?

Task parallelism

□ Why?

▣ Example: list traversal

EXAMPLE

```
void traverse_list (List l)
{
    Element e ;
    #pragma omp parallel private ( e )
        for ( e = e→first; e; e = e→next )
    #pragma omp single nowait
        process ( e ) ;
}
```

OpenMP v2.5

- Awkward!
- Poor performance
- Not composable

Task parallelism

□ Why?

▣ Example: tree traversal

EXAMPLE

```
void traverse_tree (Tree *tree)
{
    #pragma omp parallel sections
    {
        #pragma omp section
        if ( tree→left )
            traverse_tree ( tree→left );
        #pragma omp section
        if ( tree→right )
            traverse_tree ( tree→right );
    }
    process (tree);
}
```

OpenMP v2.5

- **Too many parallel regions**
 - Extra overheads
 - Extra synchronizations
 - Not always well supported

Task parallelism

- Better solution for those problems
- Main addition to OpenMP 3.0a
- Allows to parallelize irregular problems
 - ▣ unbounded loops
 - ▣ recursive algorithms
 - ▣ producer/consumer schemes
 - ▣ ...

Task parallelism

- The OpenMP tasking model
 - ▣ Creating tasks
 - ▣ Data scoping
 - ▣ Synchronizing tasks
 - ▣ Execution model

What is an OpenMP task?

- Tasks are work units which execution **may** be deferred
 - ▣ they can also be executed immediately!
- Tasks are composed of:
 - ▣ code to execute
 - ▣ data environment
 - Initialized at creation time
 - ▣ internal control variables (ICVs)

Task directive

```
#pragma omp task [ clauses ]  
    structured block
```

- Each encountering thread creates a task
 - ▣ Packages code and data environment
- Highly composable. Can be nested
 - ▣ inside parallel regions
 - ▣ inside other tasks
 - ▣ inside worksharing constructs (for, sections)

List traversal with tasks

- Why?
 - ▣ Example: **list traversal**

EXAMPLE

```
void traverse_list (List l)
{
    Element e ;

    for ( e = e→first; e; e = e→next )
    #pragma omp task
        process ( e ) ;
}
```

What is the scope of e?



Task data scoping

- Data scoping clauses
 - ▣ `shared(list)`
 - ▣ `private(list)`
 - ▣ `firstprivate(list)`
 - data is captured at creation
 - ▣ `default(shared | none)`

Task data scoping

when there are no clauses..

- If no clause
 - ▣ Implicit rules apply
 - e.g., global variables are shared
- Otherwise...
 - ▣ firstprivate
 - ▣ shared attribute is lexically inherited

List traversal with tasks

EXAMPLE

```
int a ;
void foo ( ) {
    int b , c ;
    #pragma omp parallel shared(c)
    {
        int d ;
        #pragma omp task
        {
            int e ;
            a = shared
            b = firstprivate
            c = shared
            d = firstprivate
            e = private
        } } }
```

Tip

default(none) is your friend

Use it if you do not see it clear

List traversal with tasks

EXAMPLE

```
void traverse_list (List l)
{
    Element e ;

    for ( e = e→first; e; e = e→next )
    #pragma omp task
        process ( e ) ;
}
```

e is **firstprivate**




List traversal with tasks

EXAMPLE

```
void traverse_list (List l)
{
    Element e ;

    for ( e = e→first; e; e = e→next )
    #pragma omp task
        process ( e ) ;
}
```



how we can guarantee here that the traversal is finished?

Task synchronization

- Barriers (implicit or explicit)
 - ▣ All tasks created by any thread of the current team are guaranteed to be completed at barrier exit
- Task barrier
 - #pragma omp taskwait**
 - ▣ Encountering task suspends until **child** tasks complete
 - Only direct **childs**, not descendants!

List traversal with tasks

EXAMPLE

```
void traverse_list (List l)
{
    Element e ;

    for ( e = e→first; e; e = e→next )
    #pragma omp task
        process ( e ) ;

    #pragma omp taskwait
}
```



All tasks guaranteed to be completed here

Task execution model

- Task are executed by a thread of the **team** that generated it
 - ▣ Can be executed **immediately** by the same thread that creates it
- Parallel regions in 3.0 create tasks!
 - ▣ One **implicit** task is created for each thread
 - So all task-concepts have sense inside the parallel region
- Threads can **suspend** the execution of a task and **start/resume** another

Task parallelism

□ Why?

▣ Example: list traversal

CAREFUL!

- Multiple traversal of the same list

EXAMPLE

```
List l;
```

```
#pragma omp parallel
```

```
traverse_list (l);
```

```
void traverse_list (List l)
{
    Element e ;

    for ( e = e→first; e; e = e→next )
#pragma omp task
        process ( e ) ;
}
```

Task parallelism

□ Why?

▣ Example: list traversal

EXAMPLE

```
List l;
```

```
#pragma omp parallel  
#pragma omp single  
traverse_list (l);
```

Single traversal

- One thread enters single and creates all tasks
- All the team cooperates executing them

```
void traverse_list (List l)  
{  
    Element e ;  
  
    for ( e = e→first; e; e = e→next )  
#pragma omp task  
        process ( e ) ;  
}
```

Task parallelism

- In case **task** is within a regular **counted loop** an alternative is to parallelize task creation among threads

EXAMPLE

```
/* A DIFFERENT EXAMPLE */
```

```
#pragma omp parallel
```

```
Myfunc ( );
```

Multiple traversals

- Multiple threads create tasks
- All the team cooperates executing them

```
void Myfunc ( )  
{  
    int i;  
    #pragma omp for  
        for (i=LB; i<UB; i++)  
    #pragma omp task  
        process ( i ) ;  
}
```

Task scheduling

How it works?

- Tasks are **tied** by default
 - ▣ Tied tasks are executed always by the same thread
 - ▣ Tied tasks have scheduling restrictions
 - Deterministic scheduling points (creation, synchronization, ...)
 - Another constraint to avoid deadlock problems
 - ▣ Tied tasks may run into performance problems
- Programmer can use **untied** clause to lift all restrictions
 - ▣ **Note:** Mix very **carefully** with threadprivate, critical and thread-ids

And last..

The IF clause

- If the expression of a **if** clause is **false**
 - ▣ The encountering task is suspended
 - ▣ The new task is **executed immediately**
 - with its own data environment
 - different task with respect to synchronization
 - ▣ The parent task resumes when the task finishes
 - ▣ Allows implementations to **optimize** task creation

```
#pragma omp task if (...)  
    process ( e ) ;
```



EXERCISE!