

# ASIP: Application Specific Instruction-Set Processor

## Advanced System-on-Chip Design


17.05.2016

**Michael Gautschi**

IIS-ETHZ

Luca Benini

IIS-ETHZ




# Introduction

- Contents:
  - ASIP Introduction
    - Special function unit
    - Integration in SoC
  - Commercially available tools
    - Design flow
  - Exercise session about *Tensilica Xtensa Xplorer*
    - Introduction to *Xtensa Xplorer*
    - Exercise session
- Goals:
  - Knowing the pros and cons of ASIPs
  - Learning the general concepts of ASIPs
  - Being able to design and analyze ASIPs using *Xtensa Xplorer*

# Motivation for ASIPs

- Problem:
  - GPP is too slow for some applications
  - Power consumption is not feasible for an embedded system



=> How do ASIPs achieve a better speed/energy efficiency?

# ASIP Introduction

- ASIP = **A**pplication **S**pecific **I**nstruction-Set **P**rocessor
  - Compromise between **performance and flexibility**
  - Adding **custom instruction** to the CPU as dedicated hardware
  - Designed and optimized for a **specific application**
  - Commercial tool-suites available for **fast development**
- Example where ASIPs are used:
  - Digital signal processing
  - Audio processing
  - Mobile communication: encoding/decoding/synchronization
  - Cryptographic extensions
  - And many more

# Weakness and strength of ASIPs

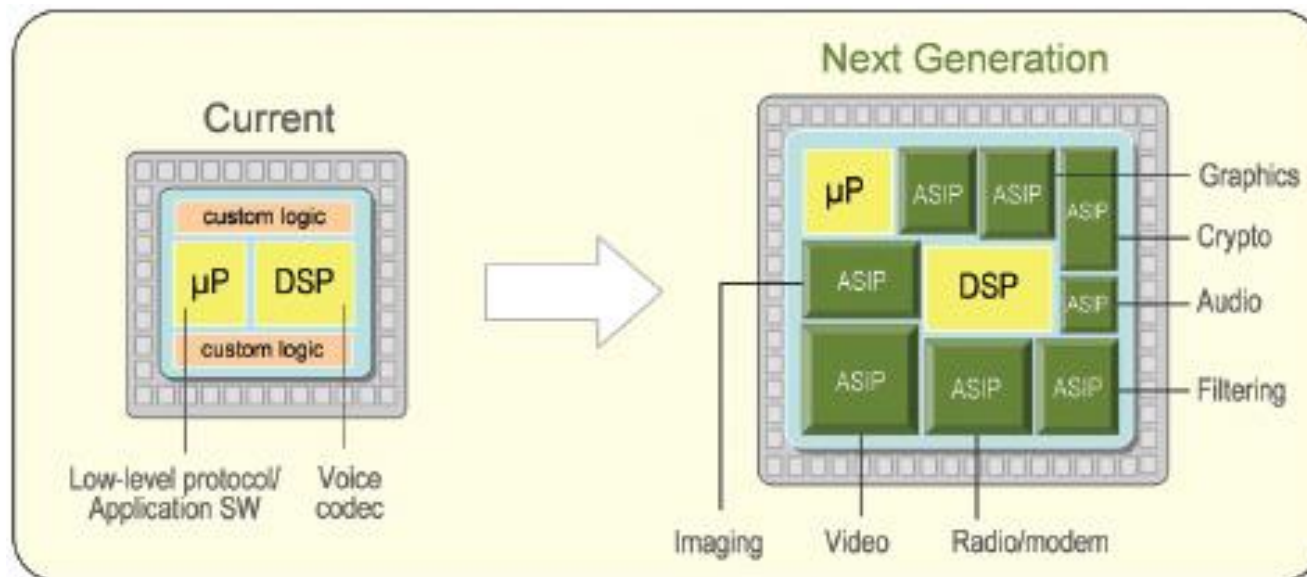
- An ASIP is a compromise between the two extremes ASIC and GPP
- ASIPs and ASICs are designed for a very specific application
  - But ASIPs are programmable and therefore almost as flexible as GPP
- ASIPs are great to be integrated in a embedded systems or SoC
  - Design effort is more on system level rather than on hardware

	<b>GPP</b>	<b>ASIP</b>	<b>ASIC</b>
<b>Performance</b>	Low	High	Very high
<b>Flexibility</b>	Excellent	Good	Poor
<b>Power</b>	Large	Medium	Small
<b>HW design</b>	Small	Large	Very large
<b>SW design</b>	Large	Large	None
<b>Reuse</b>	Excellent	Good	Poor
<b>Market</b>	Very large	Relatively large	small
<b>Cost</b>	Mainly on SW	SoC	Volume sensitive

Source: A.Nohl “*Application specific processor design: Architectures, design methods and tools*”

# ASIPs in Multicore System-on-Chip

- SoC becomes **Sea-of-Cores**
  - Many application specific processors integrated on a single chip
  - ASIPs as building blocks for heterogeneous designs

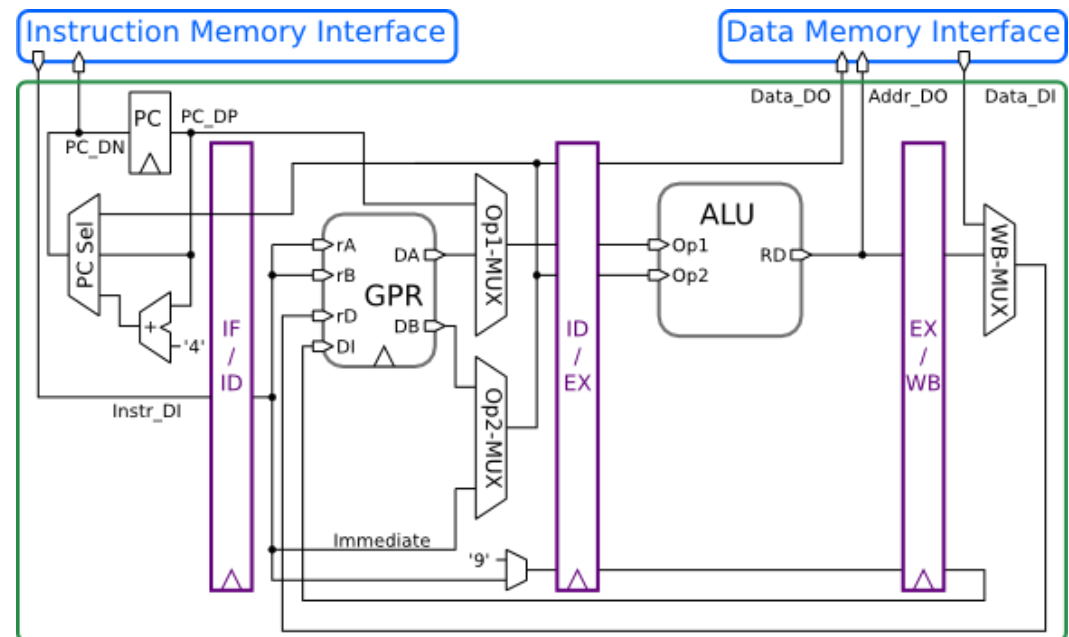


# Processor Customization

- ASIPs are customized processors with hardwired functional units in the **data path**.
- Custom instructions are executed on a special function unit (SFU)
  - Compiler maps application source code to the extended instruction set
  - New instructions can be used intrinsically in C code
- Processor data path
  - 4 stage pipeline: { **I**nstruction **F**etch, **I**nstruction **D**ecode, **E**xecute, **W**rite-**B**ack}
  - 32 bit Instruction/ data interface

Different possibilities for data path integration:

- SFU as comb. block in parallel of ALU
- Comp. unit with internal state registers
- Comp. unit with its own register file
- Shared SFUs
- Coprocessor with load/store interface

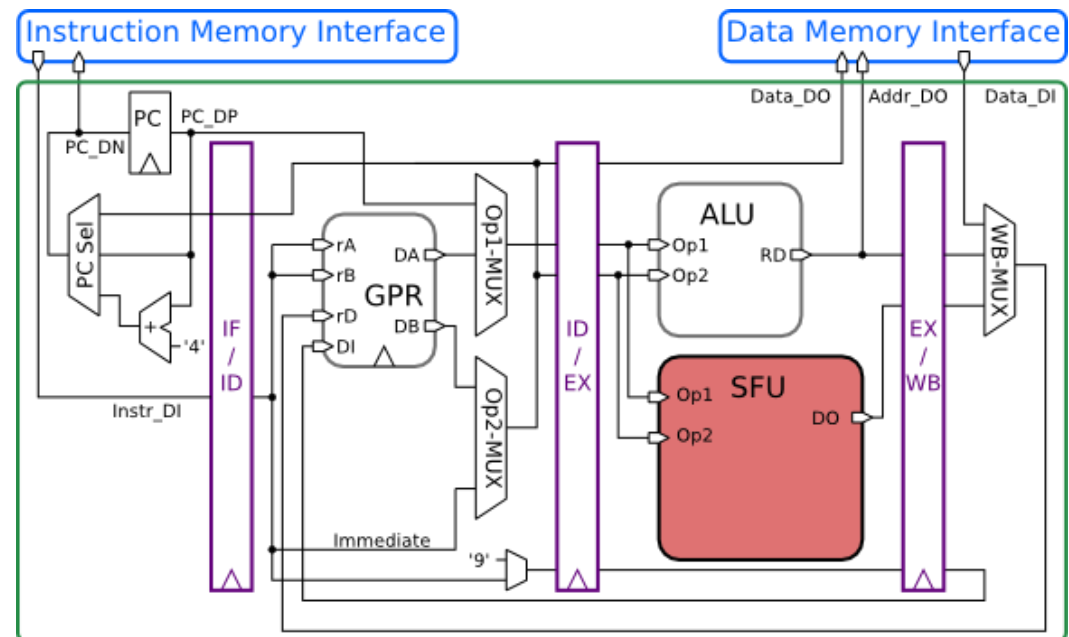


# Special Function Unit

- Combinational unit, similar to ALU
  - Gets operands from general purpose registers (2 in this example)
  - Write back the result back to GPR
  - Multi-cycle instructions to map complex functions

Example:

- Modulo operation:
  - $c = a \% b$
- Square root:
  - $c = a + \text{sqrt}(b)$



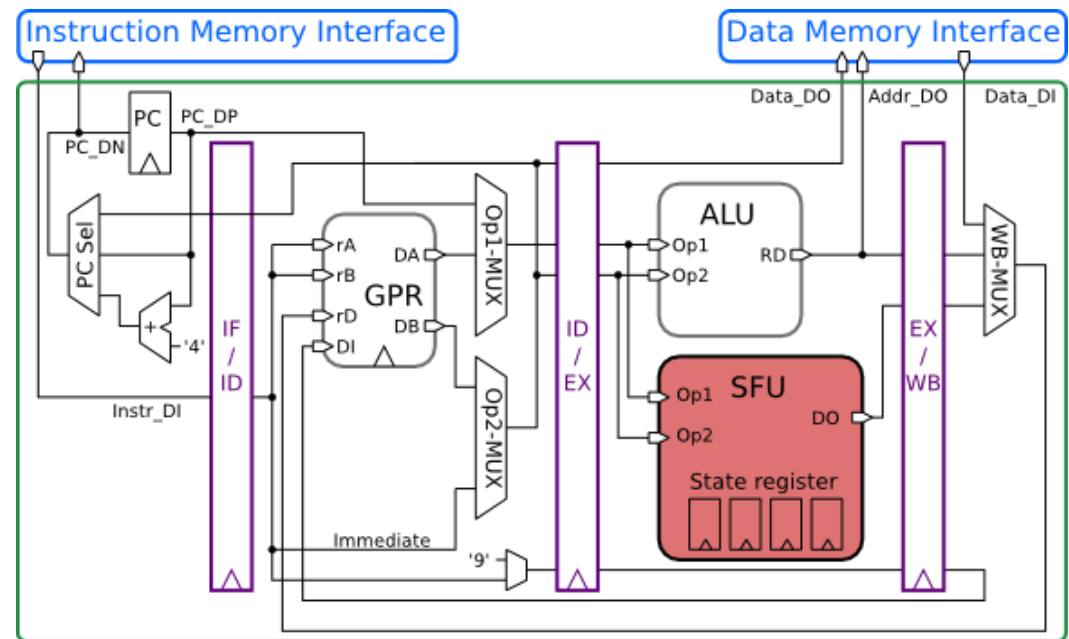


# SFU with internal state registers

- Computation unit with internal state
  - State contains content which does not need to be stored in GPR
  - Read/write state with extra read/write instructions

Example:

- Accumulator
  - $a += b[i] * c[i]$
  - $a$  stored in the state register (can be larger than 32 bit)
  - 1. initialize state to 0
  - 2. Compute MAC
  - 3. Read final state



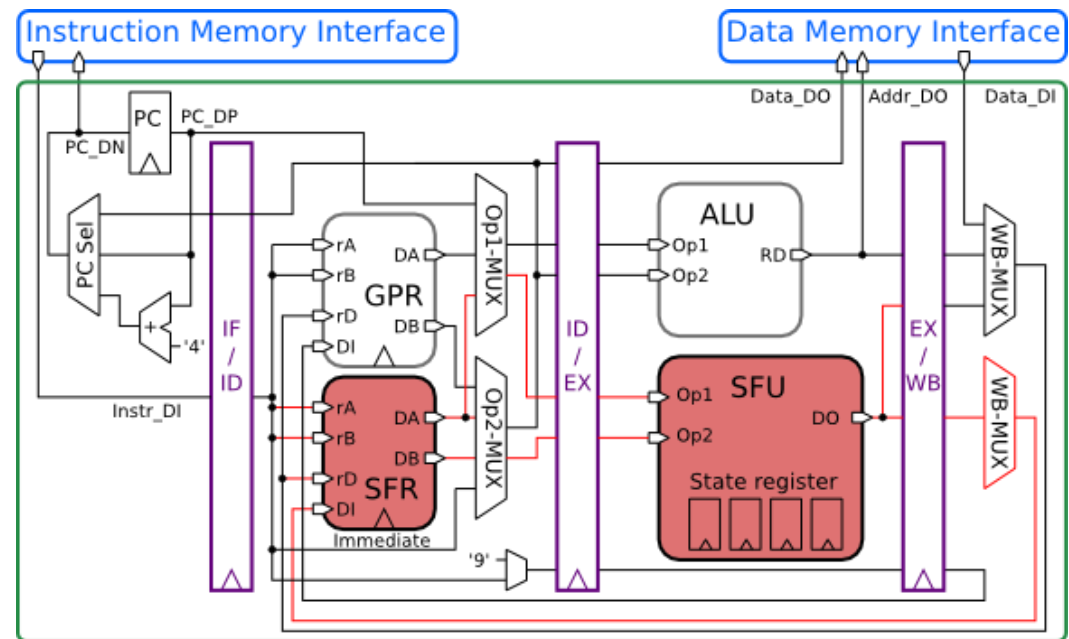
# SFU with a separate register file

- SFU reads/writes from/to the *special register file* (SPR)
  - Compiler creates move/load/store instruction to SPR
  - Custom register width optimized for the SFU
  - Custom size and number of ports at the special register file

Example:

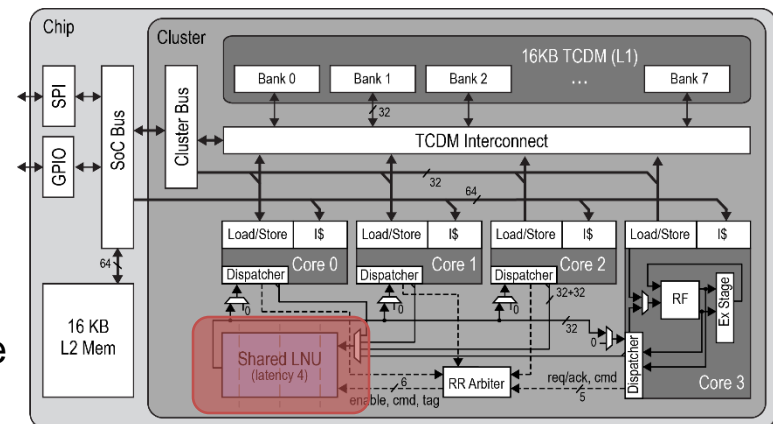
- Multiple accumulators
  - $a += b[i] * c[i]$
  - $d += f[j] * g[j]$

=> Store  $d$  and  $a$  in SPR with higher precision



# Floating-point Extensions in GPUs and Multi-core systems

- Expensive floating point operations can be computed in SFUs to approximate special functions:
  - $\sqrt{x}$ ,  $1/x$ ,  $\log_2(x)$ ,  $2^x$ ,  $\sin(x)$ ,  $\cos(x)$ , etc.
- Implemented with **function interpolators**
  - Speedup due to single cycle operation (no software emulations)
  - Limited precision (1ulp)
  - Implemented in GPUs [1,2]
- **Shared** logarithmic number unit (LNU) [3]
  - Used to interpolate floating point additions and subtractions
  - $2^x$ ,  $\log_2(x)$ ,  $\sin(x)$ ,  $\cos(x)$ ,  $\text{atan2}(y,x)$
  - Can be efficiently shared in a cluster of multiple processor cores



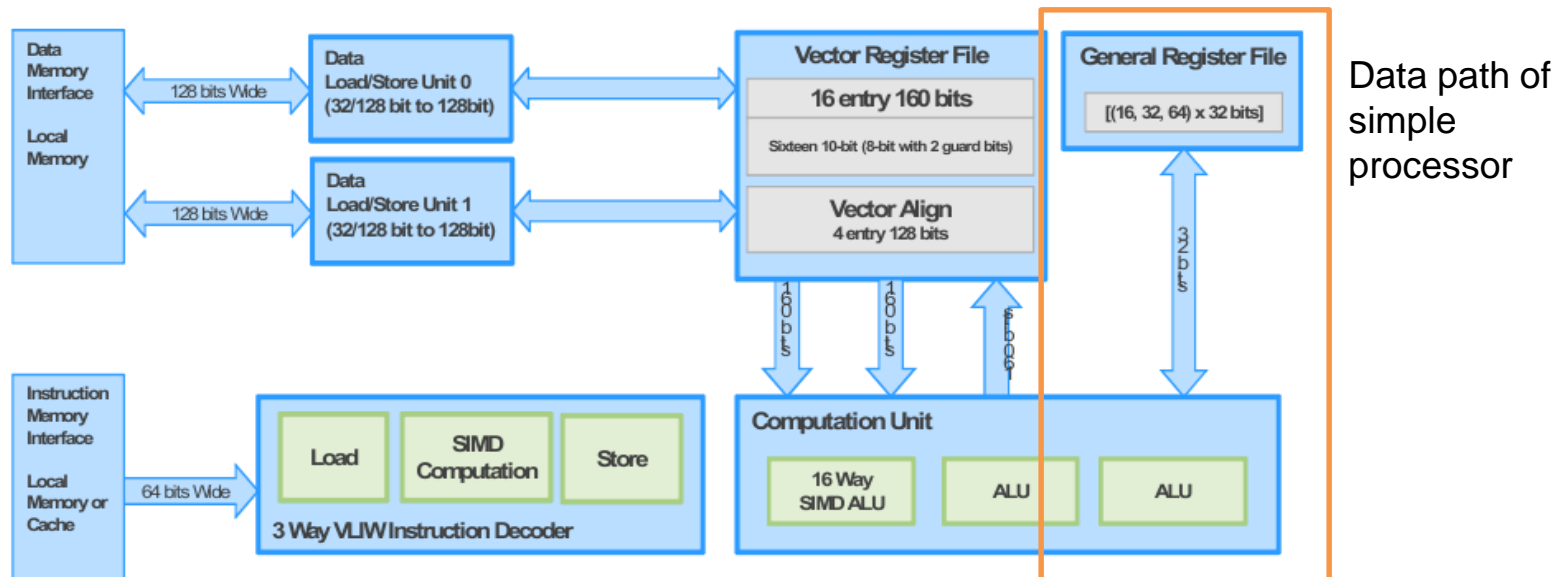
[1] Caro et. al, "High-Performance Special Function Unit for Programmable 3-D Graphics Processors", TCAS 2009

[2] Oberman et. al, "A High-Performance Area-Efficient Multifunction Interpolator", ARITH 2005

[3] Gautschi et. al, "A 65nm CMOS 6.4-to-29.2 pJ/FLOP@ 0.8 V shared logarithmic floating point unit for acceleration of nonlinear function kernels in a tightly coupled processor cluster", ISSCC 2016

# Customization with a Coprocessor


- Implement a set of instruction in a coprocessor
  - Customize the processor with more general instruction which can be used in several applications!
  - Share hardware resources for multiple instructions
  - Capable of issuing wider/more load/stores at the time
  - SIMD execution for maximum speedup
- Usage in software:
  - Program in standard C with the use of intrinsics



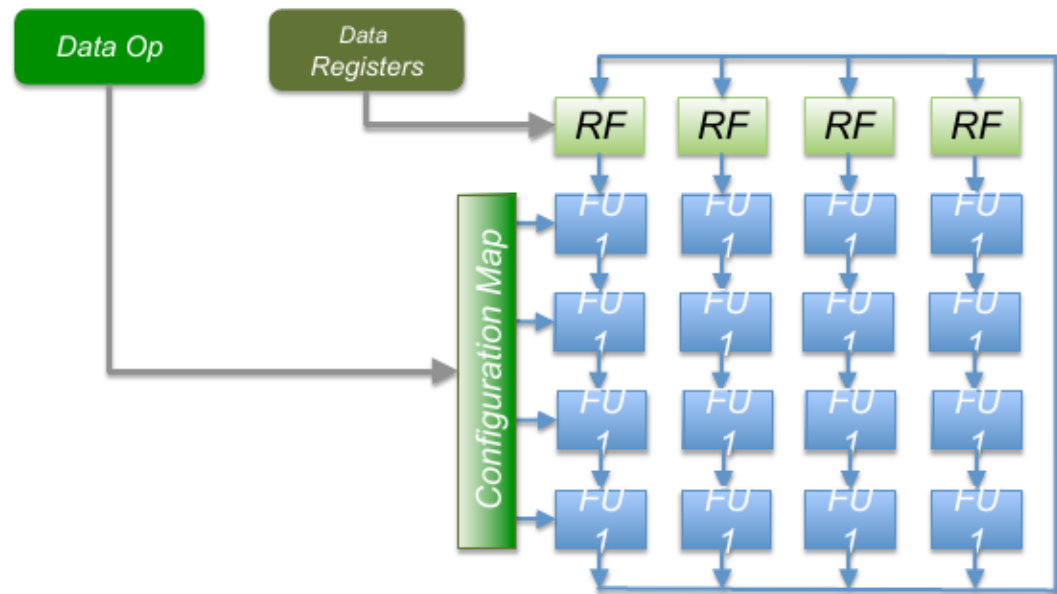
Tensilica ConnX Soft Stream Coprocessor Architecture

# Reconfigurable ASIPs

- Design SFU for more than one single application!



- Replace SFU by configurable array of simple functions
- Configure the array to process the desired function
- FPGA-like approach
  - Efficient if not necessary to reconfigure very often



Source: Berekovic, TU-Braunschweig 2010

# Commercially available ASIP/VLIW design tools

- LISA (Synopsys)



- Xtensa  
(Tensilica/Cadence)

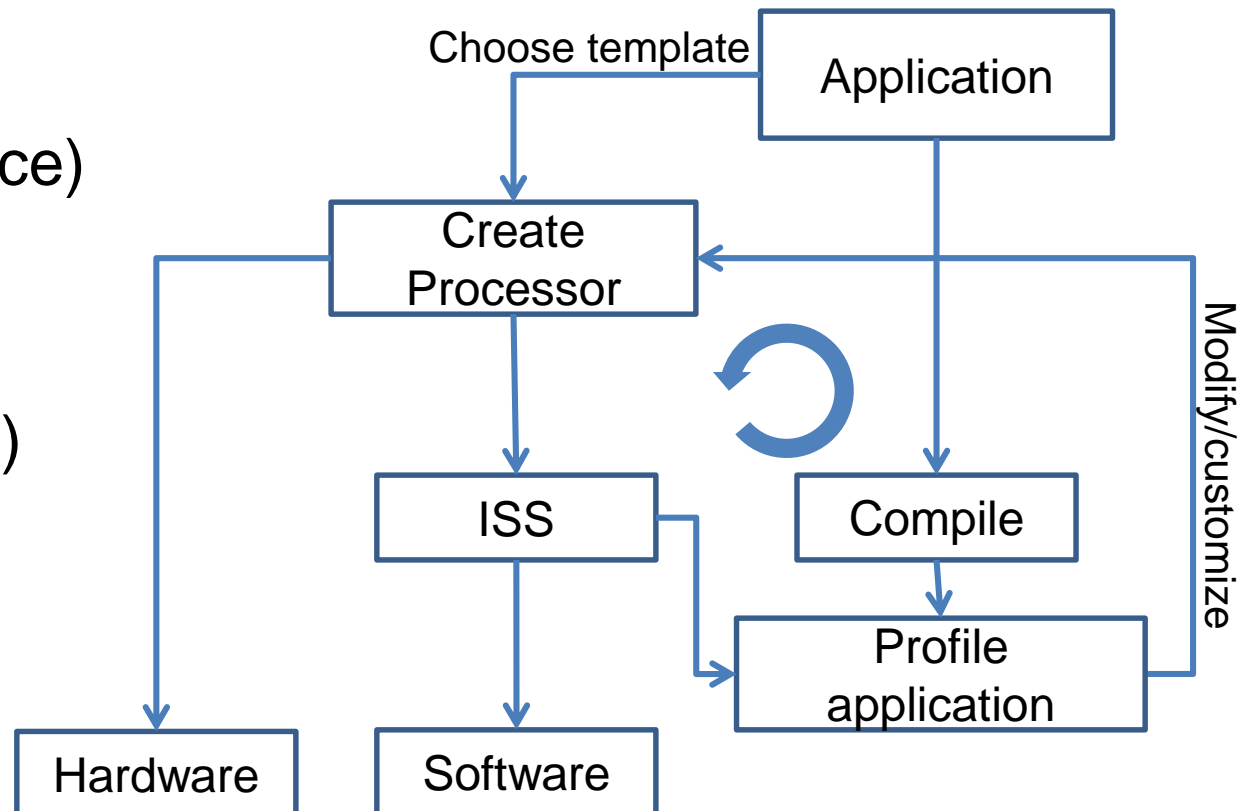


- OptimoDE (ARM)

- Codasip



## Design Flow:

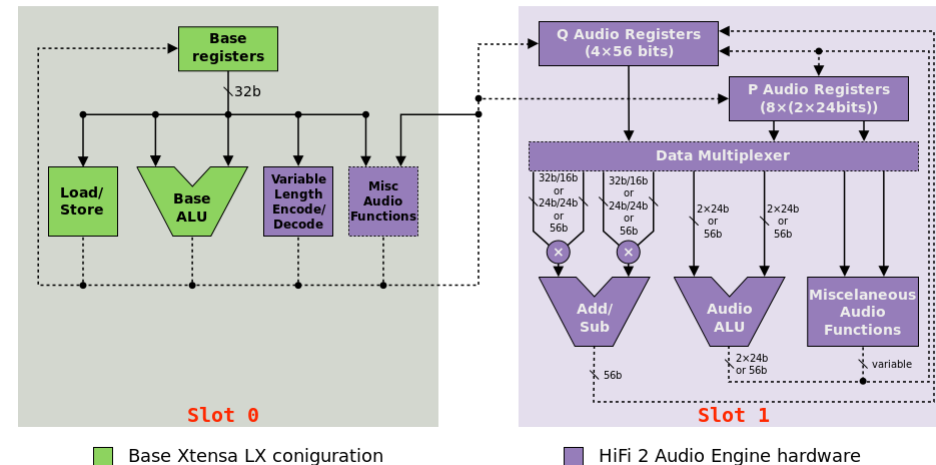


# Where is it used?

- **AMD Radeon R9-290**
  - High-end graphics card with TrueAudio
  - TrueAudio is a co-processor based on the HIFI core architecture of Cadence
  - AMD supports audio processing via TrueAudio in all its new chips (architecture generation GCN 1.1)



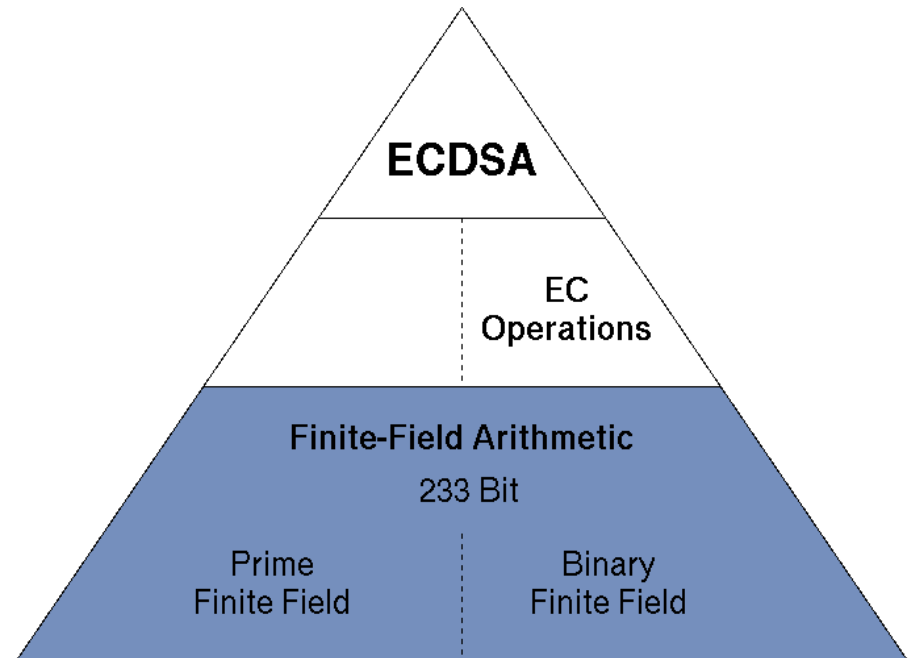
Datapath and registers of Cadence Tensilica's HiFi 2



- **Samsung Galaxy S3**
  - Voice recognition used by S-Voice
  - Uses low power HIFI xtensa core

# Cryptographic Example: ECDSA signature verification

- **E**lliptic **C**urve **D**igital **S**ignature **A**lgorithm (ECDSA)
  - Signature validation with **E**lliptic **C**urves (EC)
- ECDSA requires prime and binary finite field operations
- Add, sub, mult, div, etc. defined in the finite-field arithmetic
  - We used the NIST B-233 curve
- Flexibility is good!
- Possible to support different algorithms and different standards with one hardware!



- Finite field operations are not suitable for general purpose architectures  
=> hardware much more efficient!



# Cryptographic Example: Base processor configuration

- Processor configuration:
  - 5 stage integer pipeline
  - 32 32bit general purpose registers
  - 16x16 bit multiplier
  - 2 KB I\$, D\$
- Estimated performance (65nm LP technology):
  - Max speed: 344 MHz (worst case)
  - Area: 83 kGE
  - Power : 20 mW (not verified)
- ECDSA algorithm ported from a former semester project [1]

[1] Semester Thesis by A. Traber, S. Stucki, 2014, A Unified Multiplier Based Hardware Architecture for Elliptic Curve Cryptography

# ECDSA Signature Verification Algorithm

Profiling results using Xtensa Xplorer:

Operation	Total (cycles)	Total (%)	# function calls	Code size (bytes)
ECDSA Verification	46'674'997			6518
$GF(2^{233})$ multiplication	42'063'643	90.1%	2'309	366
16x16 bit binary f. field mult.	34'199'189	73.3%	591'104	157
$GF(2^{233})$ squaring	2'706'852	5.7%	2'472	446
others	1'904'502	21.0%	-	4'930

- Simple squaring operation on a 233 bit binary field:
  - Insert '0' bit between each input bit
    - e.g. '1101' => '01010001'
  - Reduce resulting 466 bit number to 233 bit
- Requires masking and shifting in C -> not efficient at all!
- Hardware architectures can do such operations in one cycle!  
=> add custom instructions for 16bit mult, and squaring

# Optimizing the multiplication in $GF(2^{233})$

## Original C code

```
// shortened C-code: binary field 16 bit
// multiplication

uint16_t a, b;    // input
uint32_t prod = 0; // output

for(i = 0; i < 16; i++) {
    if((a & (1 << i))
        prod ^= b << i;
}

~60 cycles
```

## New C code

```
// intrinsic function call
prod = BinMul16(a,b);

// a, b, prod are stored
```

Speedup of  
factor 60!

## “TIE” – Tensilica Instruction Extensions

```
// shortened BinMul operation
operation BinMul16 {out AR res, in AR a, in AR b}{}{
    wire [31:0] temp0 = (a & (1 << 0)) ? (b << 0)          : 32'b0
    wire [31:0] temp1 = (a & (1 << 1)) ? (temp0^(b << 1))   : temp0;
    wire [31:0] temp2 = (a & (1 << 2)) ? (temp1^(b << 2))   : temp1;
    // ...
    wire [31:0] temp15 = (a & (1 << 15)) ? (temp14^(b << 15)) : temp14;
    assign res = temp15;
}
```

1 cycle only!

# Results: Performance

- Profiling results of specialized circuit
- Reduced code size: 5.1 KB vs 5.7 KB

Operation	Total (cycles)	# cycles before	Speedup
ECDSA Verification		46'674'997	
$GF(2^{233})$ multiplication		42'063'643	
16x16 bit binary f. field multiplication		34'199'189	
$GF(2^{233})$ squaring		2'706'852	
others		1'904'502	

- Comparison to HW-architecture:[2]
  - Coprocessor with 16 bit datapath requires ~1'850'000 cycles (12 kGE)
  - Factor 3.3 slower

[2] M. Gautschi, M. Mühlberghuber et.al., SIRIOUS: A tightly coupled ECC Coprocessor for the OpenRISC

# Results: Area and Timing

- Synthesis results with UMC 65 nm technology

Core comparison	Max Speed (MHz)			Area (kGE)			
	Synthesis		Estimation	Synthesis			Estimation
	TC	WC		TC	WC	344 MHz	
Basic configuration	857*	580*	344	103	107	98	83
With hardware extensions	850*	589*	344	107	111	100	85
Instruction extension				4.1	4.2	2.3	2.2

\* Synthesized without memories; expected drop by 200-400MHz when synthesizing with a 2KB cache

- Conclusion:
  - Flexible architecture fully integrated in an application processor
  - 7.5x speedup at very low costs and design time (< 1 week)
  - Only 2.2 kGE hardware overhead, (datapath of co-processor 12 kGE)

# Summary

- ASIPs are widely used in application specific domains
- ASIP can increase energy efficiency
- Possible to create complex SoCs with ASIPs
- ECDSA example is based on a mini-project completed by Sven Stucki

# Q&A

